

GKI

Ending kernel fragmentations on Android

Juhung Park

✉️ arter97@dgist.ac.kr

GitHub arter97

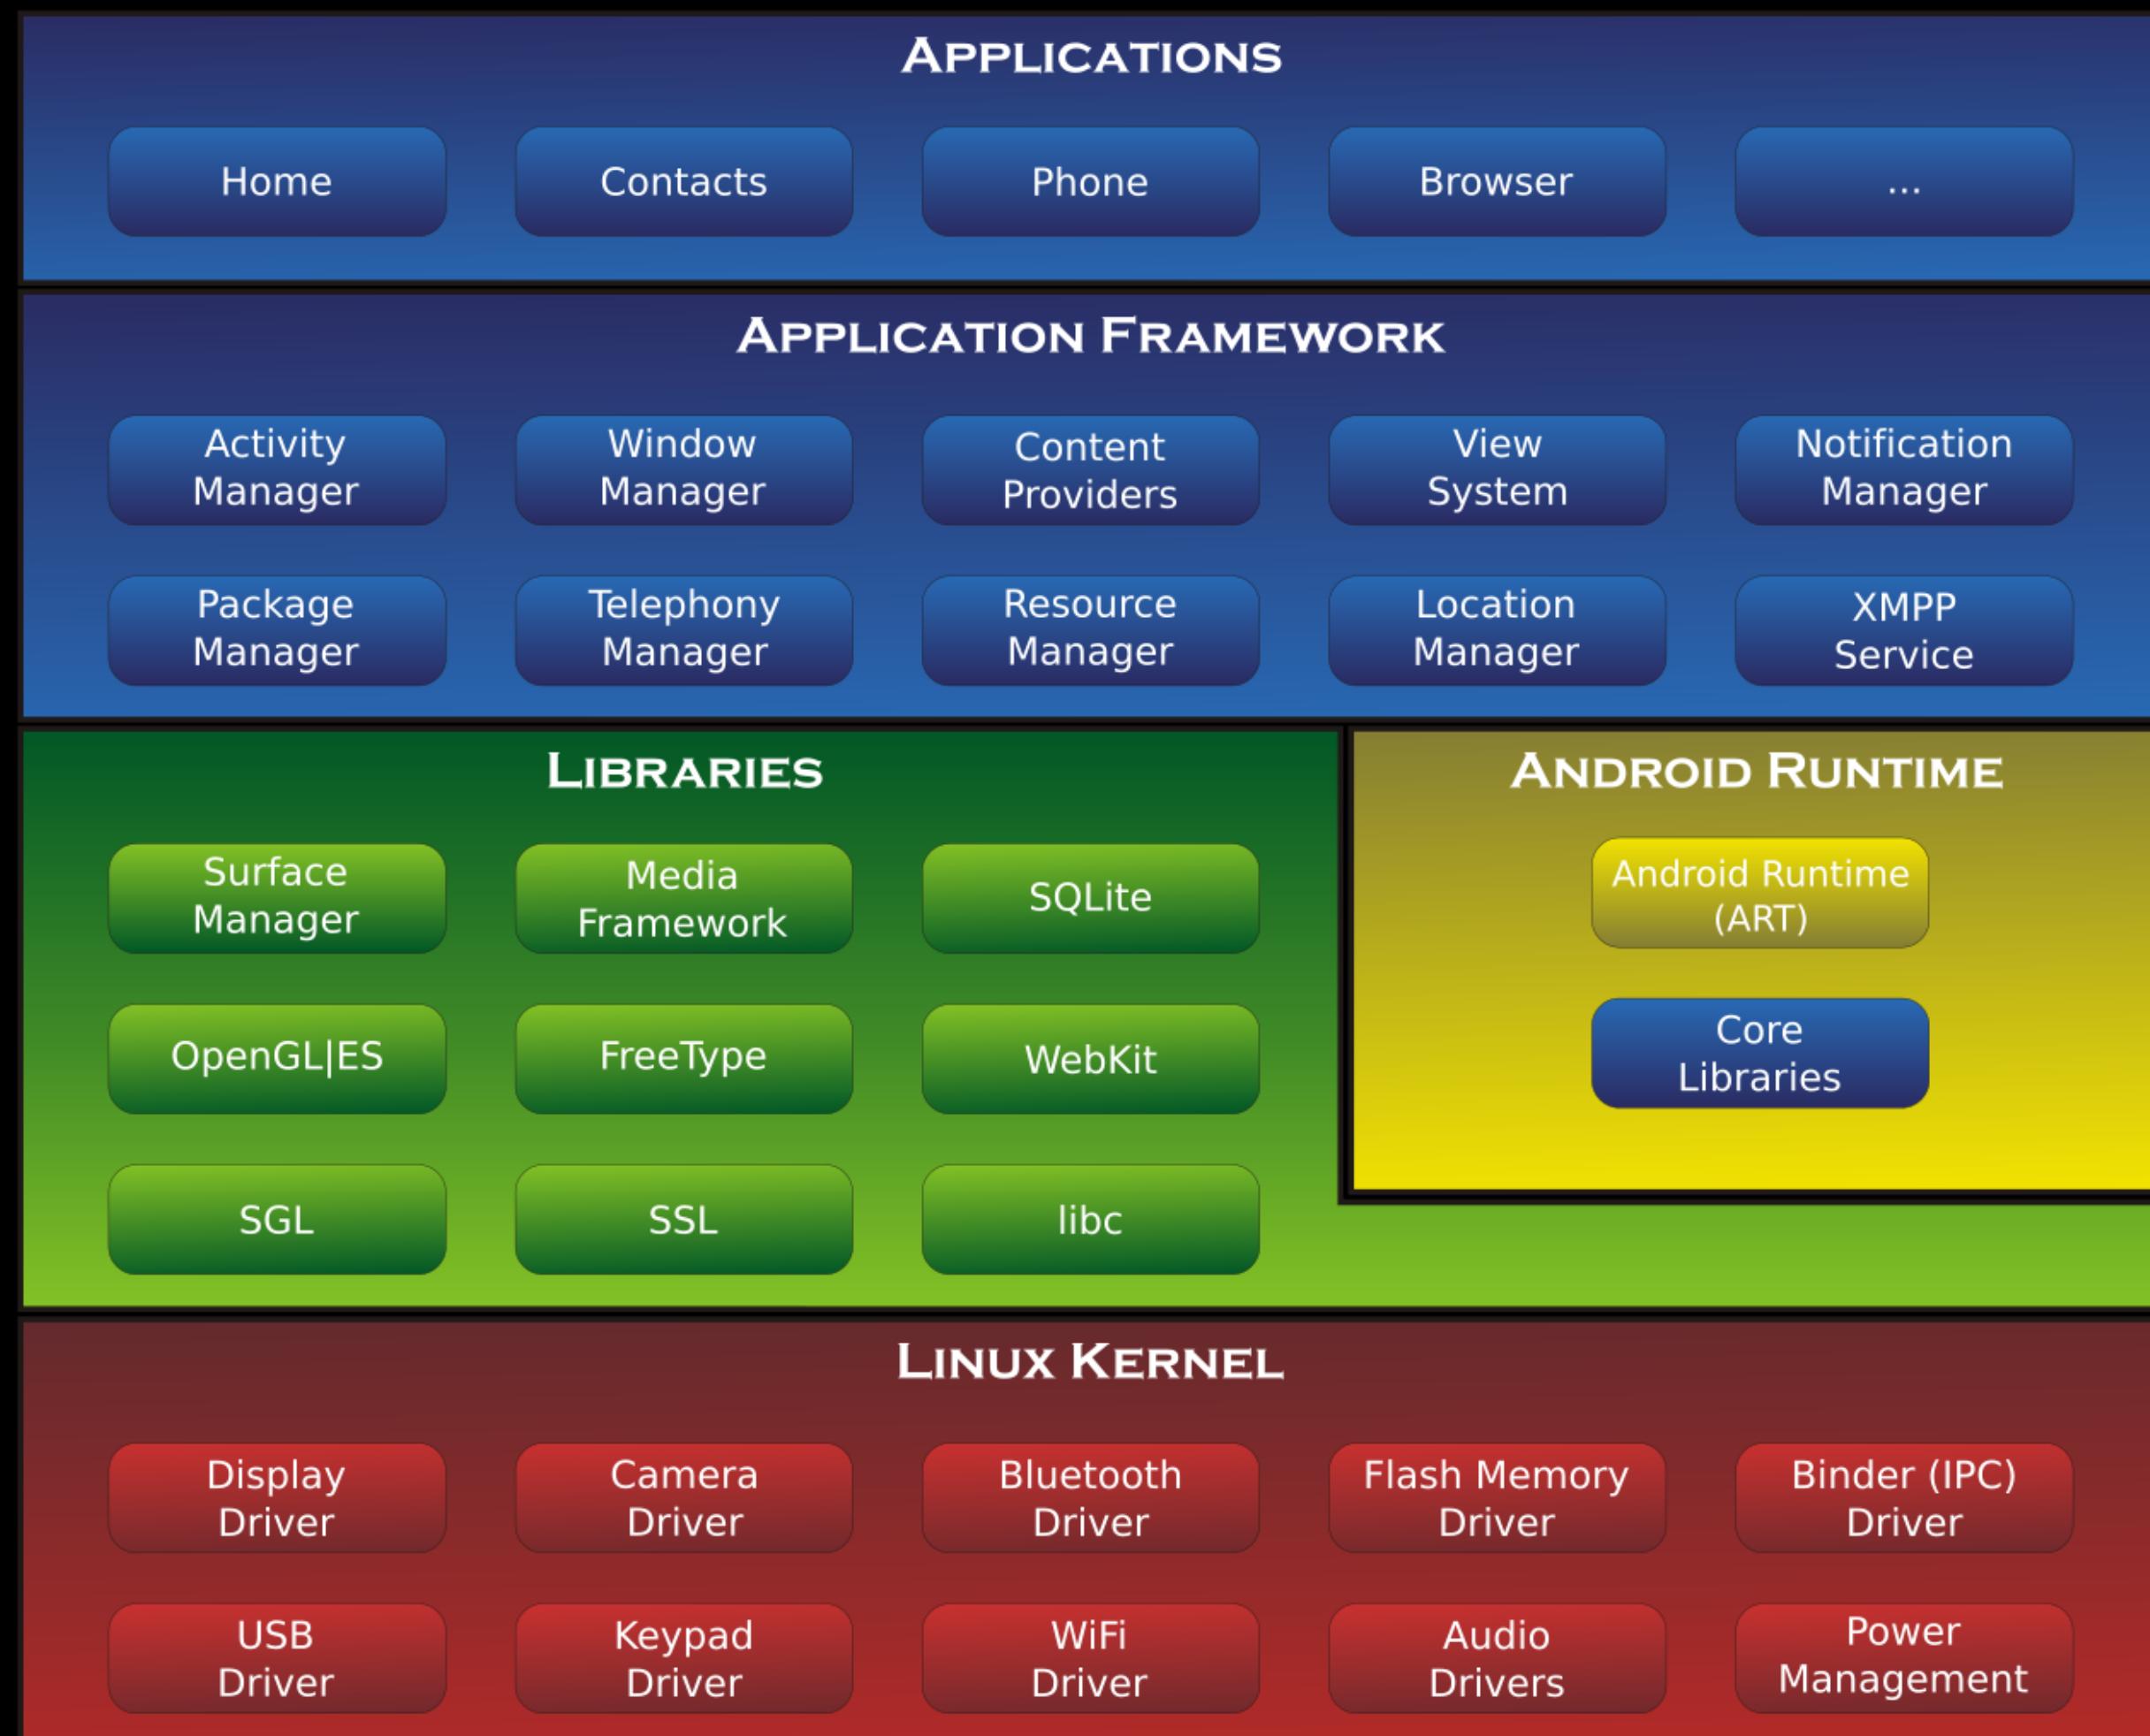
Twitter @arter97

About me

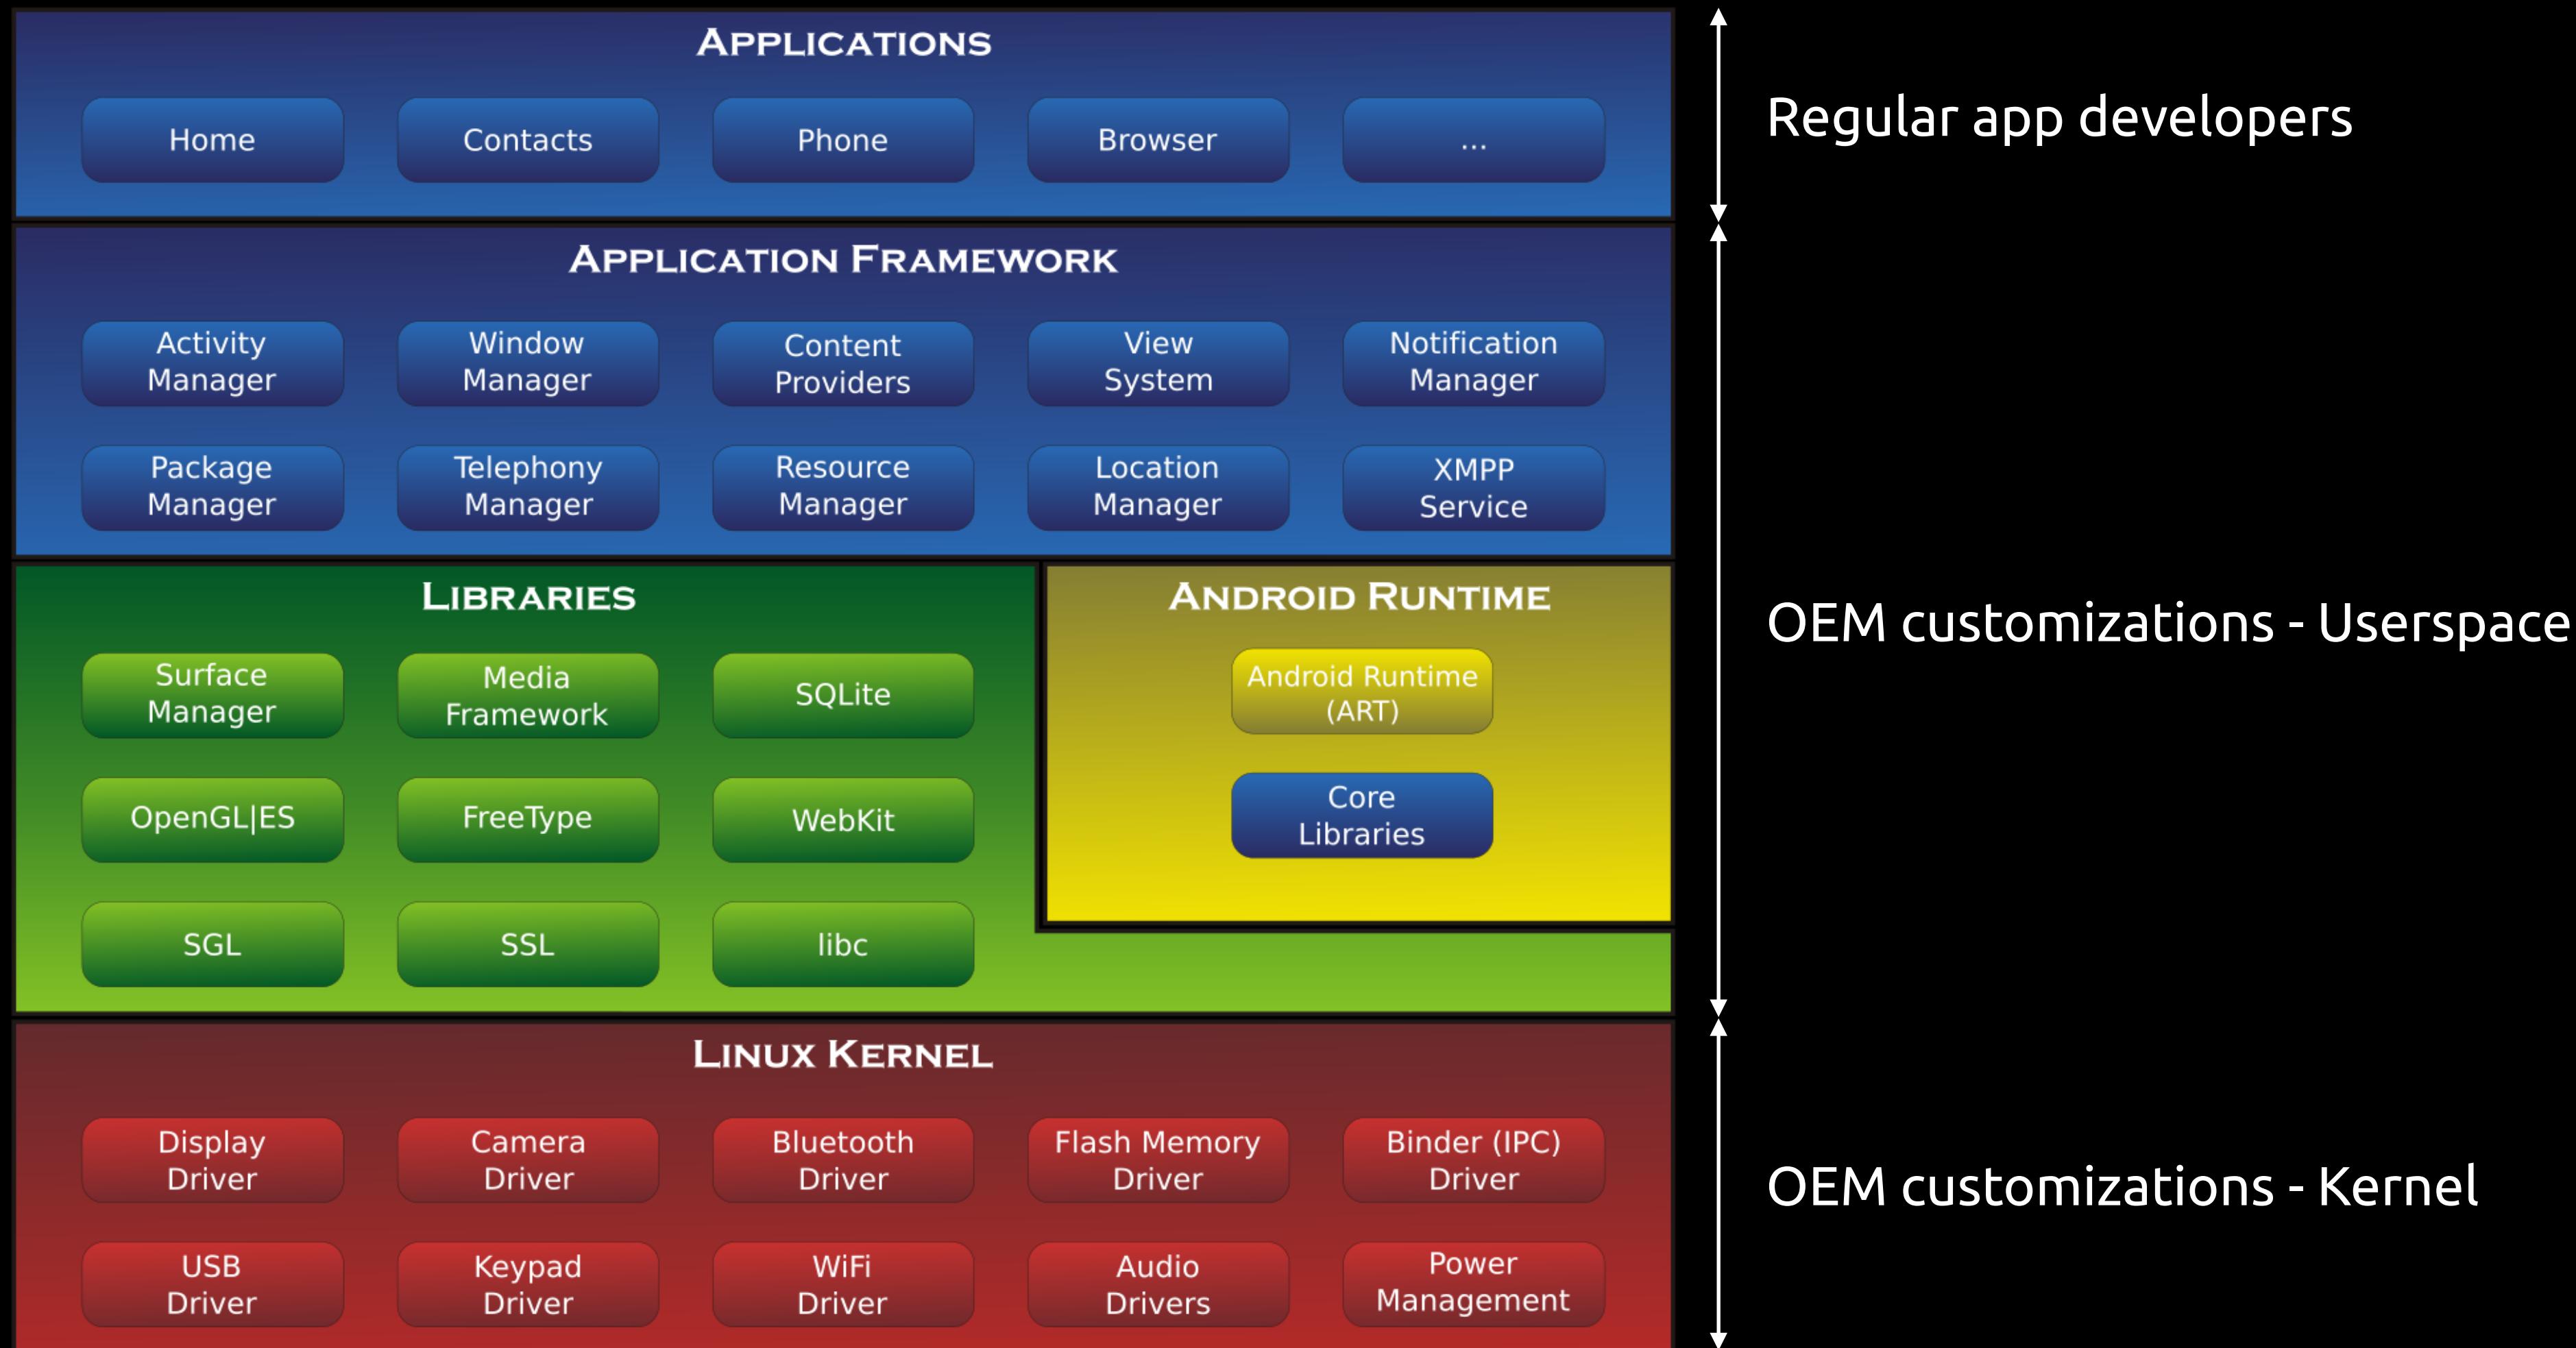
- Hobbyist Android / kernel developer since 2012
- Ph.D. student @ DGIST DataLab
- Interests: storage, Android
- 5th kernel meet-up (2019)
 - Historical overview of Android and its kernel



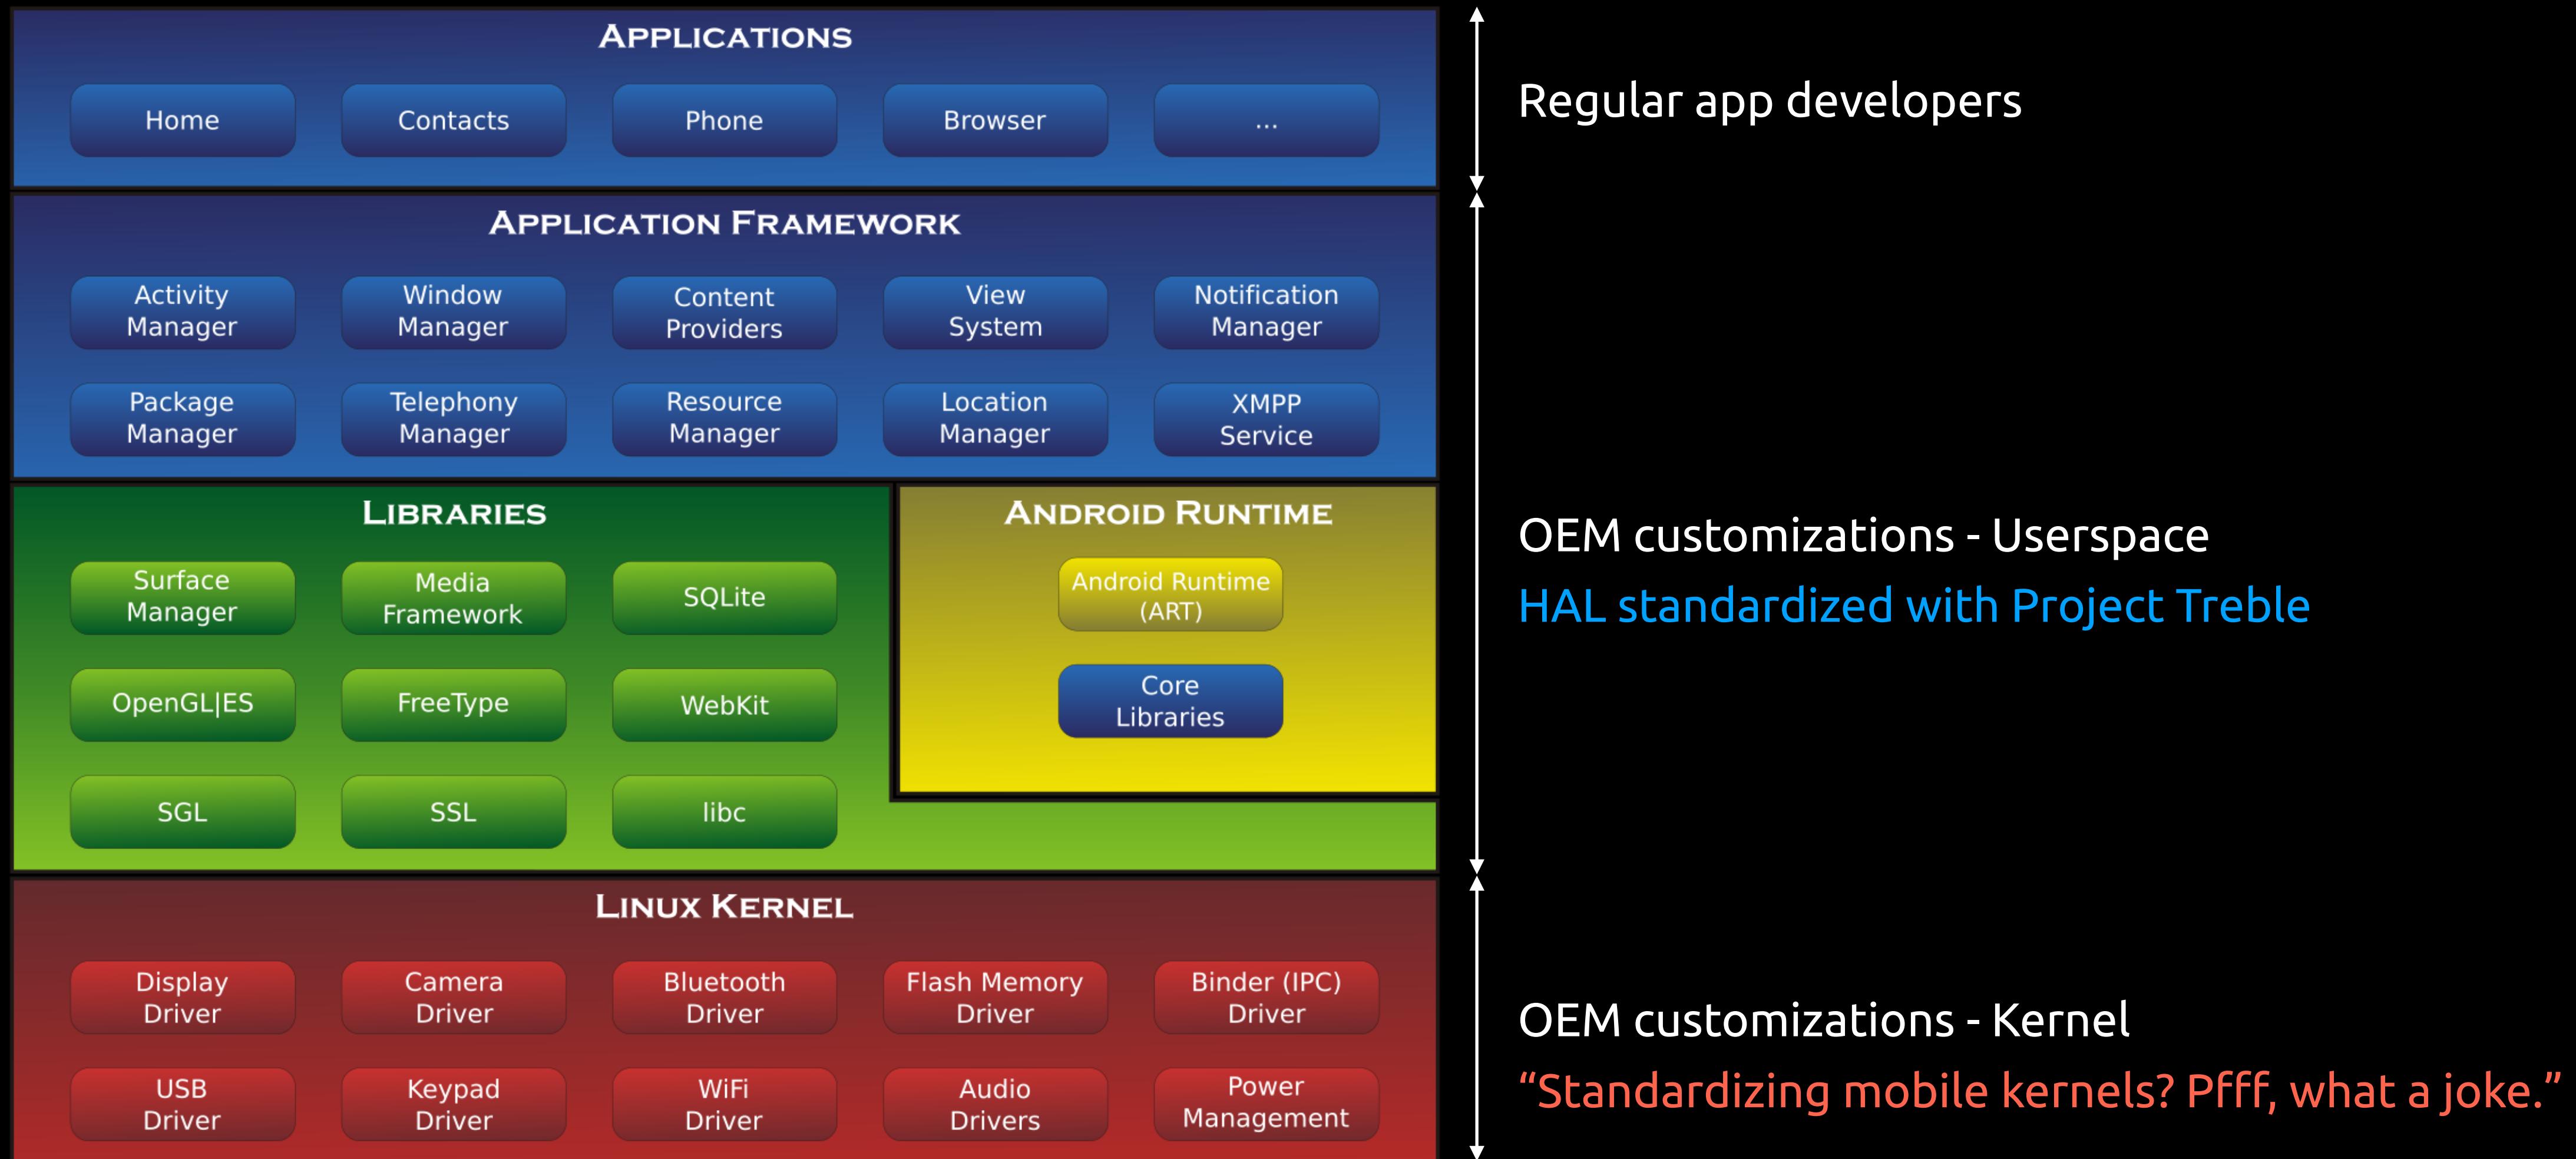
Android fragmentation



Android fragmentation



Android fragmentation



ARM kernel fragmentation

```
From      Linus Torvalds <>
Date      Thu, 17 Mar 2011 19:50:36 -0700
Subject   Re: [GIT PULL] omap changes for v2.6.39 merge window
```

Gaah. Guys, this whole ARM thing is a f*cking pain in the ass.

You need to stop stepping on each others toes.

<...>

Somebody needs to get a grip in the ARM community. I do want to do these merges, just to see how screwed up things are, but guys, this is just ridiculous. The pure amount of crazy churn is annoying in itself, but when I then get these "independent" pull requests from four different people, and they touch the same files, that indicates that something is wrong.

Somebody in the ARM community really needs to step up and tell people to stop dicking around.

* arm64 is slightly better with "device tree"

```
$ tree -L 1 arch/arm
arch/arm
├── boot
├── common
├── configs
├── crypto
├── include
├── Kconfig
├── Kconfig.debug
├── Kconfig-nommu
├── kernel
├── kvm
├── lib
├── mach-actions
└── mach-alpine
...
├── mach-vt8500
├── mach-zx
├── mach-zynq
└── Makefile
├── mm
├── net
├── nwfpe
├── oprofile
├── plat-omap
├── plat-orion
├── plat-pxa
├── plat-samsung
├── plat-versatile
├── probes
└── tools
├── vdso
└── vfp
└── xen
```

94 directories, 4 files

Mainline kernel?

The screenshot shows a news article from Phoronix. The header features the Phoronix logo in white on a dark background. Below the header is a green navigation bar with links: ARTICLES & REVIEWS, NEWS ARCHIVE, FORUMS, PREMIUM, CONTACT, and CATEGORIES. The main content area has a white background. The article title is "Qualcomm Snapdragon 8 Gen 1 Sees Timely Support With The Mainline Linux 5.17 Kernel". Below the title is the author information: "Written by Michael Larabel in Arm on 12 January 2022 at 06:54 AM EST. 3 Comments". To the left of the text is the Arm logo, which consists of a blue square with the word "arm" in white. The text discusses Qualcomm's timely mainline support for their latest SoCs, comparing it favorably to other vendors' slower adoption.

Qualcomm only announced the Snapdragon 8 Gen 1 and X65 platforms at the end of November but already they have managed to provide timely mainline support for these latest high-end SoCs. This is great to see compared to the days of slow to materialize mainline support for new Arm SoCs, which still persists among some vendors with either belated mainline support or only focusing on vendor downstream kernels. The big batch of Arm SoC/platform changes have landed for [Linux 5.17](#).

<https://www.phoronix.com/news/Linux-5.17-Arm>

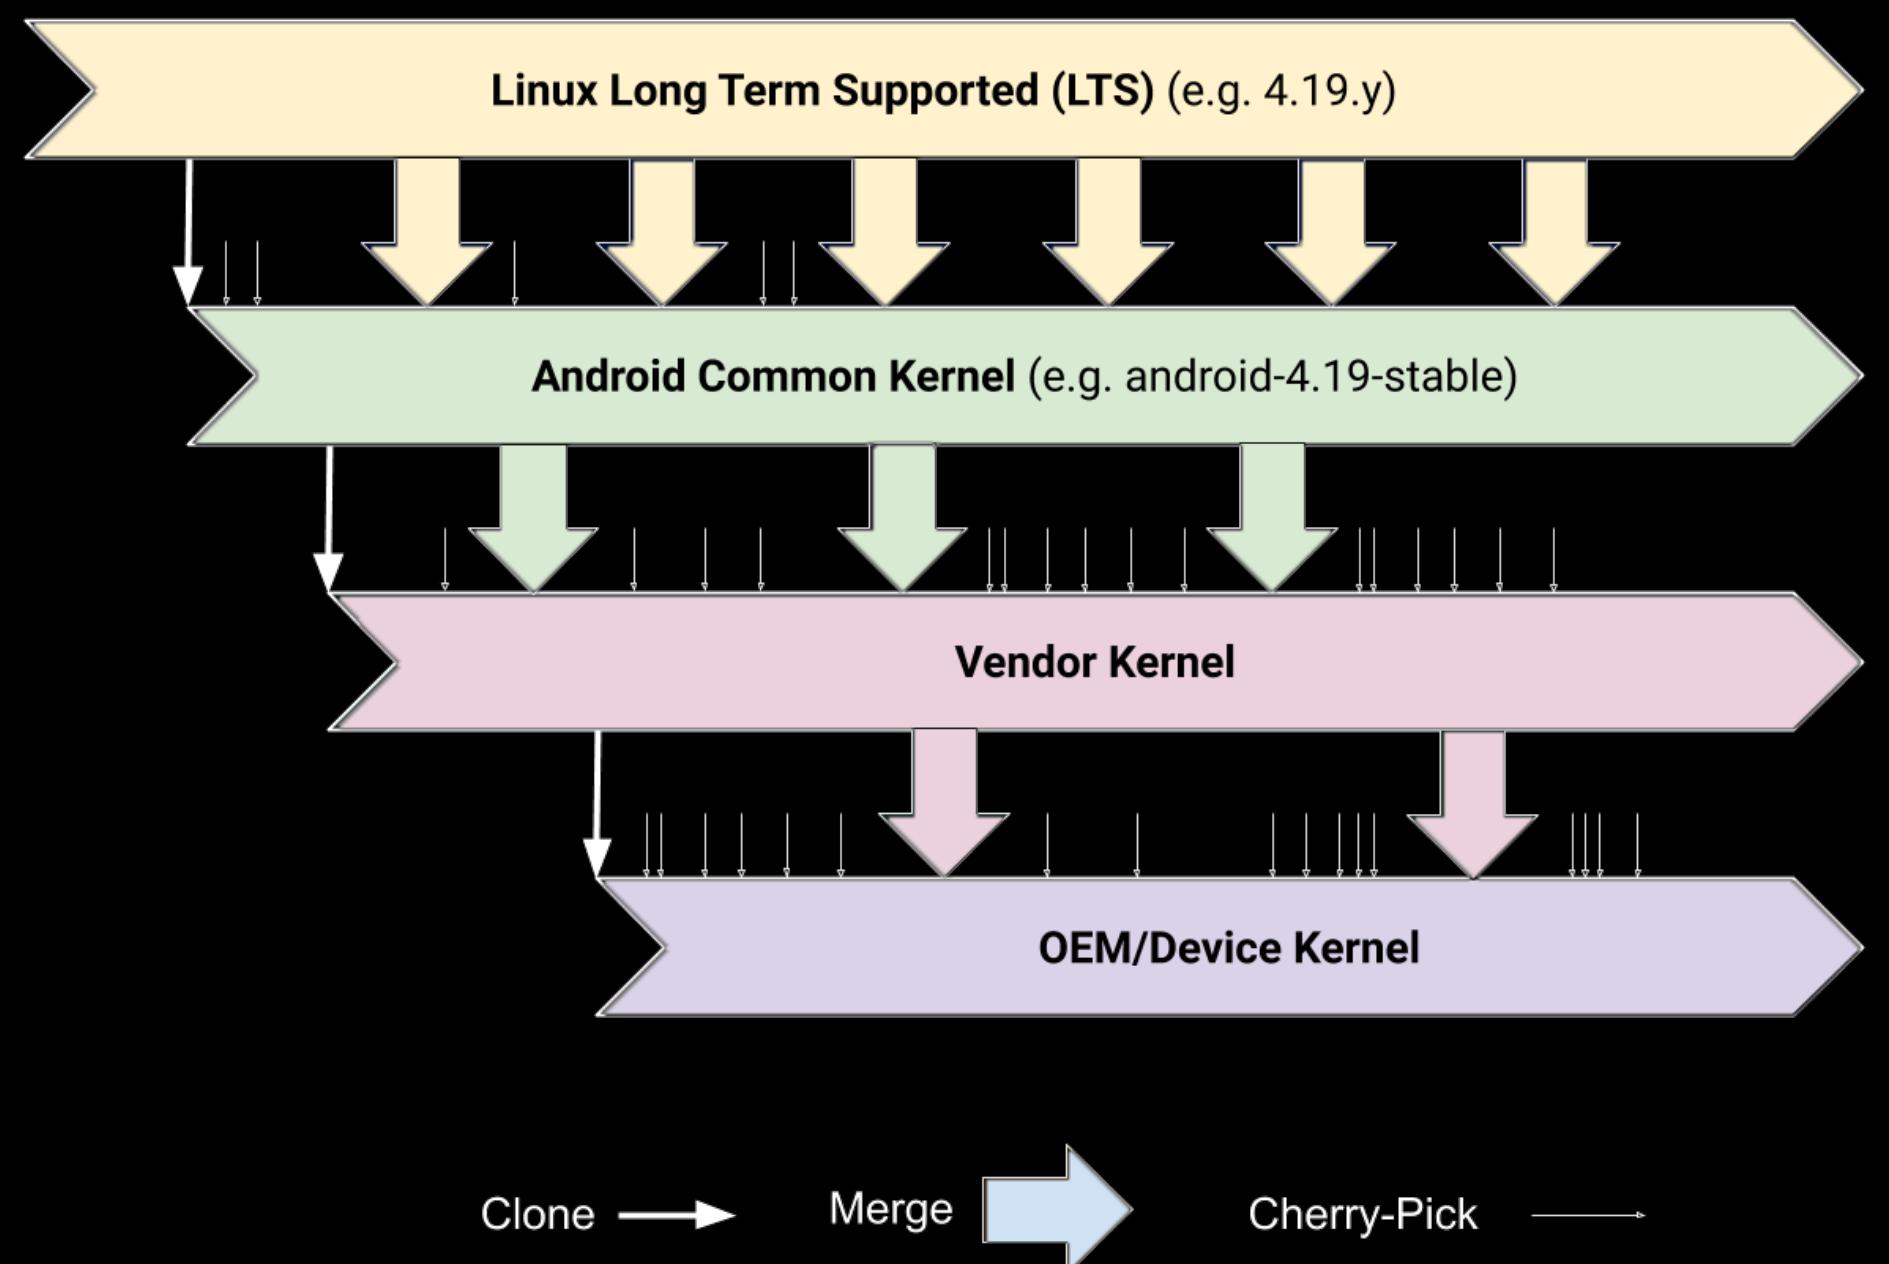
Mainline kernel?

Nope.

	Mainline	Downstream
CPU	✓	✓
GPU	▲	✓
Storage	▲	✓
Modem	✗	✓
Wi-Fi	✗	✓
Bluetooth	✗	✓
Camera	✗	✓
Battery	✗	✓
Audio	✗	✓
Biometrics	✗	✓
Touch panel	✗	✓

Android downstream kernel

- Upstream: The Linux kernel from kernel.org
- AOSP: Additional Android-specific patches from AOSP common kernels
- Vendor: SoC and peripheral enablement and optimization patches from vendors
- OEM/device: Additional device drivers and customizations



Android downstream kernel

Cost of fragmentation

- Security updates are labor intensive
- Merging linux-stable is difficult
 - 90% of kernel security issues are resolved from linux-stable
- Android platform upgrade conflicts
 - New version of Android may require new kernel changes
- Difficult to upstream changes

GKI

Generic Kernel Image

- Unifying the core kernel and separating SoC / board support
- Maintaining a stable ABI (KMI) for vendor modules
 - Allows modules and the core kernel to be updated independently
 - Huh?

ABI stability in the Linux kernel?

stable-api-nonsense

Greg Kroah-Hartman <greg@kroah.com>

This is being written to try to explain why Linux **does not have a binary kernel interface, nor does it have a stable kernel interface**.

You think you want a stable kernel interface, but you really do not, and you don't even know it.

<https://www.kernel.org/doc/Documentation/process/stable-api-nonsense.rst>

ABI stability in the Linux kernel?

stable-api-nonsense

Ever wondered why your NVIDIA / VMware drivers are broken or rebuilt upon kernel upgrades?

<https://www.kernel.org/doc/Documentation/process/stable-api-nonsense.rst>

ABI stability in the Linux kernel?

stable-api-nonsense

- The compiler may change data structure order and padding
- Different kernel configuration may produce incompatible changes
 - Structures containing different fields
 - Unimplemented features
 - Memory alignment changes
- New commits and fixes may change data structure

<https://www.kernel.org/doc/Documentation/process/stable-api-nonsense.rst>

ABI stability in the Linux kernel? What's the upstream's answer?

“Simple, get your kernel driver into the main kernel tree.”

Infeasible for Android OEMs for obvious reasons.

Maintaining a stable KMI

Kernel Module Interface

- Only a single configuration, gki_defconfig, is used to build the core kernel
- KMI is maintained within the same LTS and Android version of a kernel
 - e.g., android13-5.10, android12-5.10, android13-5.15
- Only one **specific LLVM toolchain** can be used to build the kernel and modules
- Symbols used from modules are whitelisted/exported
 - EXPORT_SYMBOL() isn't enough
 - Vendors have to upload patches to change the symbols list

Maintaining a stable KMI

KMI breakages

```
diff --git a/include/linux/mm_types.h b/include/linux/mm_types.h
index 5ed8f6292a53..f2ecb34c7645 100644
--- a/include/linux/mm_types.h
+++ b/include/linux/mm_types.h
@@ -339,6 +339,7 @@ struct core_state {
 struct kioctx_table;
 struct mm_struct {
 struct {
+    int dummy;
    struct vm_area_struct *mmap;           /* list of VMAs */
    struct rb_root mm_rb;
    u64 vmacache_seqnum;                  /* per-thread vmacache */
```

Maintaining a stable KMI

Detecting KMI breakages

Leaf changes summary: 1 artifact changed

Changed leaf types summary: 1 leaf type changed

Removed/Changed/Added functions summary: 0 Removed, 0 Changed, 0 Added function

Removed/Changed/Added variables summary: 0 Removed, 0 Changed, 0 Added variable

'struct mm_struct at mm_types.h:372:1' changed:

 type size changed from 6848 to 6912 (in bits)

 there are data member changes:

[...]

Maintaining a stable KMI KABI struct reservation

- Paddings are reserved for future uses
- Mostly taken from RedHat's RHEL/CentOS KABI patches

[Code](#) [Pull requests](#) [Actions](#) [Security](#) [Insights](#)**ANDROID: GKI: add some padding to some driver core structures**

These structures are fundamental to implementing fw_devlink and sync_state(). Since they are still evolving, add some padding in case we need to backport any important bug fixes.

```
struct device_link  
struct class  
struct fwnode_handle  
struct fwnode_link
```

Bug: 183615740

Signed-off-by: Saravana Kannan <saravanak@google.com>

Signed-off-by: Greg Kroah-Hartman <gregkh@google.com>

Change-Id: Id9daf7cf9ae5d94fb0134144f8220a241ccbaf8

↳ android13-5.15 + android13-5.15-2022-06 + android13-5.15-2022-07 + android13-5.15-2022-08 + android13-5.15-2022-09 + android13-5.15-arcvm + android13-5.15-lts

Saravana Kannan authored and Carlos Llamas committed on Jun 19

Showing 3 changed files with 8 additions and 0 deletions.

Filter changed files

include/linux

device.h

device

class.h

fwnode.h

2 include/linux/device.h

```
@@ -612,6 +612,8 @@ struct device_link {  
 612    612      struct kref kref;  
 613    613      struct work_struct rm_work;  
 614    614      bool supplier_preactivated; /* Owned by consumer probe. */  
 615    +      ANDROID_KABI_RESERVE(1);  
 616    +      ANDROID_KABI_RESERVE(2);  
 615    617    };  
 616    618  
 617    619    static inline struct device *kobj_to_dev(struct kobject *kobj)  
.....
```

Maintaining a stable KMI

Downstream kernel development

- If data structures are changed, send it to Android Common Kernel (ACK)
 - Preferably before KMI is frozen
- Refactor so that data structure doesn't need to diverge from ACK
- Paddings are used here too

What about un-upstreamable changes?

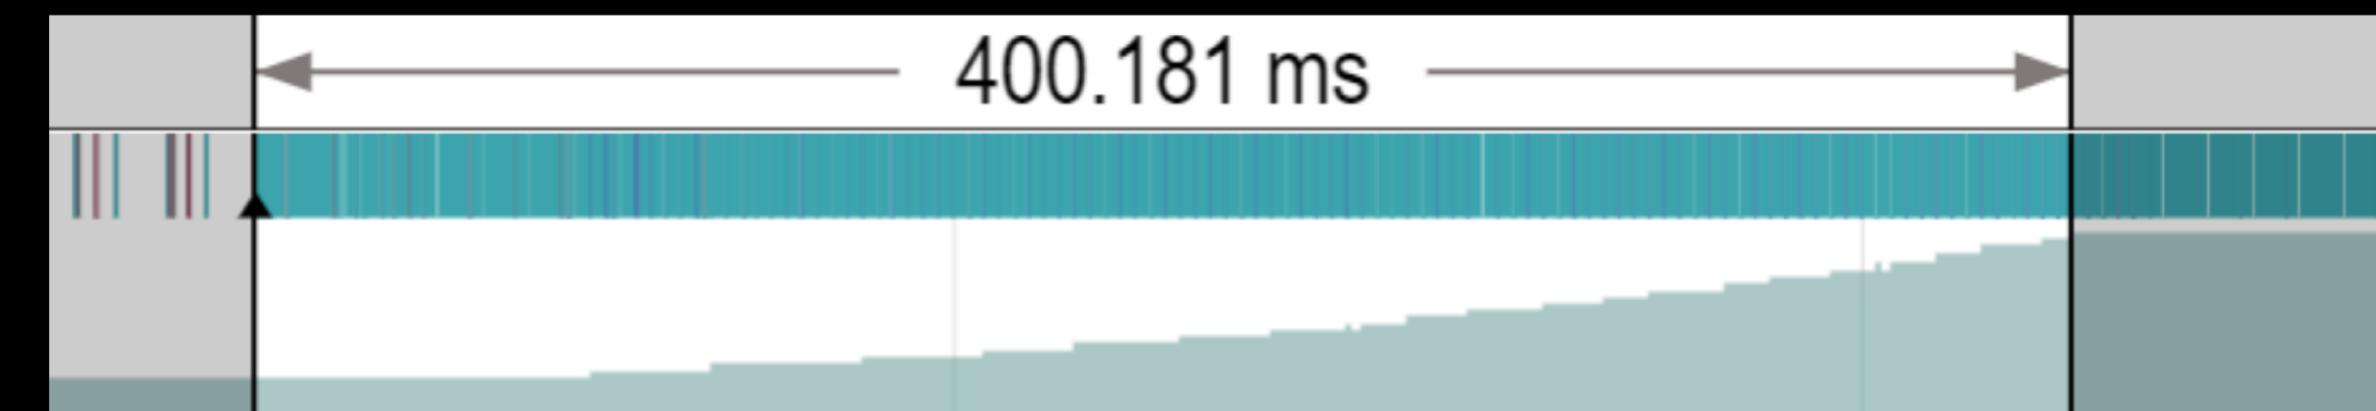
- Vendor-hooks to modify default behavior of GKI
 - Preferred than “weak” functions
- Not enough; some OEMs heavily modifies the kernel
 - Very core - arch, mm, sched, etc
 - Some are even hw-specific
- GKI 2.0’s goal - “Don’t introduce significant performance or power regressions when replacing the product kernel with the GKI kernel.”

Downstream kernel development

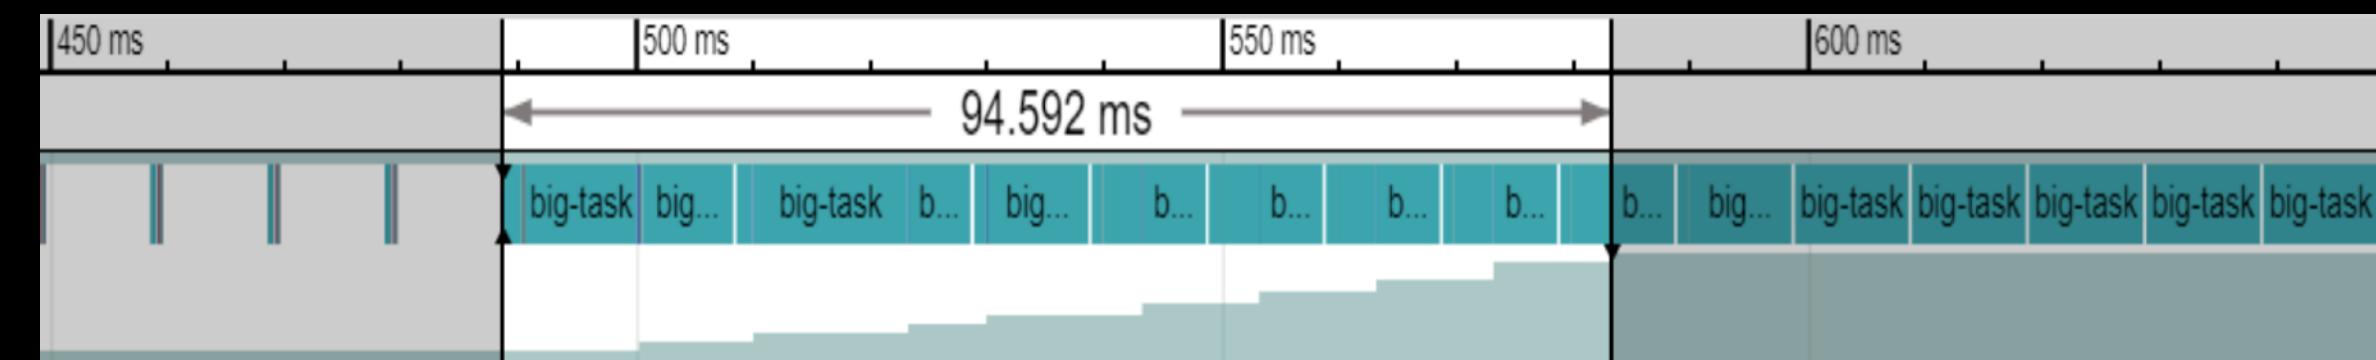
Case-study: Qualcomm's WALT scheduler

- PELT (Per Entity Load Tracking) uses moving load average to calculate the next task demand

- Slow ramp-up



- WALT (Window Assisted Load Tracking) references previous N windows to predict the next task demand



Downstream kernel development

Case-study: Qualcomm's WALT scheduler

- How many *core* files were changed with v4.19?

- Excluding walt-specific files

```
git grep -i walt LA.UM.9.12.1.r1-03100-SMxx50.QSSI12.0 -- kernel/sched | grep -vP 'walt.*\.\?:' | wc -l  
250
```

- What about v5.10? (GKI 2.0)

```
git grep -i walt google/android13-5.10 -- kernel/sched | grep -vP 'walt.*\.\?:' | wc -l  
0
```

- How?

Downstream kernel development

Case-study: Qualcomm's WALT scheduler

- Just a bunch of “hooks”

```
$ git log --oneline --no-merges google/android12-5.4..google/android13-5.10 -- kernel/sched | grep ANDROID: | grep -i vendor | grep -i hook  
19e41a340441b ANDROID: sched: Add vendor hook for cpu distribution functions  
65735b81dd25c ANDROID: sched: Add vendor hook for update_rq_clock_pelt  
3f9db3f711817 ANDROID: sched: Add vendor hook for rt util update  
e3356ca0a6c24 ANDROID: sched: Add vendor hook for util-update related functions  
0975fd934e9e3 ANDROID: Add vendor hook for the sugov_get_util  
cf33d6fae045e ANDROID: vendor_hooks: add a hook to control the delay time of frequency up and down.  
61b66296b2a24 ANDROID: Add vendor hook for the uclamp_util  
478f3ee1d3e5b ANDROID: Add vendor hook to the deadline scheduler  
f9fcdaeb7006 ANDROID: sched: remove regular vendor hooks for 32bit execve  
13af062abf9b3 ANDROID: vendor_hooks: Export the tracepoints sched_stat_sleep and sched_waking to let module probe them  
27c285003d2a8 ANDROID: sched: Add vendor hook to select ilb cpu  
a6bb1af39d11e ANDROID: vendor_hooks: Export the tracepoints sched_stat_iowait, sched_stat_blocked, sched_stat_wait to let modules probe them  
fe580539f6cec ANDROID: vendor_hooks: Add hooks for account irqtime process tick  
f3f8d55011837 ANDROID: sched: Add vendor hooks for update_load_avg  
54f66141a8834 ANDROID: sched: Add vendor hooks for sched.  
58b10706f3c8d ANDROID: sched: Add vendor hooks to compute new cpu freq.  
50aa353c6957b ANDROID: sched: Add vendor hooks for cpu affinity.  
b79d1815c400c ANDROID: psi: Add vendor hooks for PSI tracing  
63399b4e2e25b ANDROID: vendor_hooks: Add hooks for account process tick  
187306ab1aef ANDROID: Partial revert of 06881e01b564 ("ANDROID: sched: Add vendor hooks for override sugov behavior")  
b542f4c389e9a Revert "ANDROID: sched: Add vendor hooks for skipping sugov update"  
475aea007d656 ANDROID: sched: Add vendor hook for util_est_update  
53e8099784435 ANDROID: vendor_hooks: Add hooks for scheduler  
cc1f93cb20c46 ANDROID: sched: Add vendor hook for uclamp_eff_get  
067bf187ef4c0 ANDROID: sched: move vendor hook to check scheduling nice value  
4b88cf85247f6 ANDROID: sched: Add vendor hook for uclamp_eff_value  
97368fc2dcc29 ANDROID: Add a vendor hook that allow a module to modify the wake flag  
f5998fbf2dbbc ANDROID: sched: Add vendor hook for cpu_overutilized  
a56f081c5b2ec ANDROID: sched: Add restricted vendor hooks in CFS scheduler  
18ebdc37464a1 ANDROID: sched: add vendor hooks for bad scheduling  
4b5f345c7391b ANDROID: sched: Add restrict vendor hooks for balance_rt()  
18b294fa4c884 ANDROID: sched: add em_cpu_energy vendor hook  
11d9d07f3f686 ANDROID: schedutil: add vendor hook for adjusting util to freq calculation  
fb03f7bef8ea4 ANDROID: sched: Export symbol for vendor RT hook function  
95fd09fa3a0f1 ANDROID: sched: Export symbol for vendor RT hook function  
d31ea0950a30a ANDROID: sched/fair: fix place_entity() vendor hook  
292f430816f36 ANDROID: Sched: Add restricted vendor hooks for scheduler  
f34f38632fb72 ANDROID: sched/core: Add vendor hook to change task affinity  
846bf8e8cb392 ANDROID: sched: Add vendor hooks for skipping sugov update  
06881e01b564a ANDROID: sched: Add vendor hooks for override sugov behavior  
147a9b3d9eab7 ANDROID: sched: Add vendor hooks for find_energy_efficient_cpu  
878495dacd22d ANDROID: sched: Add restrict vendor hooks for load balance  
ebce8ec6bd55e ANDROID: sched: rt: rearrange invocation of find_lowest_rq() vendor hook  
a2ca8408de618 ANDROID: sched: add restrict vendor hook to modify load balance behavior  
2935d588b1378 ANDROID: sched: Use normal vendor hook in scheduler tick  
9ad8ff902e39f ANDROID: vendor_hooks: add waiting information for blocked tasks  
f39f3ac200dbb ANDROID: sched: add vendor hooks to handle scheduling priority  
16330b35600f0 ANDROID: Add vendor hooks to the scheduler  
a1fc1fba460be ANDROID: sched: add restrict vendor hook to modify task placement policy in EAS
```



aosp-mirror / kernel_common

Public mirror

mirrored from <https://android.googlesource.com/kernel/common.git>

 Notifications

 Fork 520

 Star 434



 Code  Pull requests  Actions  Security  Insights

ANDROID: sched: gki: add padding to some structs to support WALT

[Browse files](#)

Add padding to below structs to support WALT based accounting:

1. struct cpu_topology
2. struct task_struct
3. struct sched_domain_shared
4. struct task_group
5. struct root_domain
6. struct rq

To accommodate potential future changes, reserving more memory than what WALT needs today.

Bug: 171858786

Change-Id: If6d901174fc7963be3ae44daa799cb2953669ec1

Signed-off-by: Satya Durga Srinivasu Prabhala <satyap@codeaurora.org>

 Satya Durga Srinivasu Prabhala authored and [toddkj](#) committed on Oct 30, 2020

1 parent 8d5dc2c

commit 7a2a316228bcd80821a15fed6cd70c80a5daa182

 Showing 4 changed files with 14 additions and 1 deletion.

Split

Unified

Downstream kernel development

Downstream kernel development

Case-study: Qualcomm's WALT scheduler - Vendor hooks for altering behaviors

```
$ git grep -i android_ KERNEL.PLATFORM.1.0.r1-11200-kernel.0 kernel/sched/walt | grep rvh
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/fixup.c:static void android_rvh_show_max_freq(void *unused, struct cpufreq_policy *policy,
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/fixup.c:register_trace_android_rvh_show_max_freq(android_rvh_show_max_freq, NULL);
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/preemptirq_long.c:register_trace_android_rvh_irqs_disable(note_irq_disable, NULL);
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/preemptirq_long.c:register_trace_android_rvh_irqs_enable(test_irq_disable_long, NULL);
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/preemptirq_long.c:register_trace_android_rvh_preempt_disable(note_preempt_disable, NULL);
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/preemptirq_long.c:register_trace_android_rvh_preempt_enable(test_preempt_disable_long,
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_cpu_cgroup_online(void *unused, struct cgroup_subsys_state *css)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_cpu_cgroup_attach(void *unused,
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_wake_up_new_task(void *unused, struct task_struct *new)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_update_cpu_capacity(void *unused, int cpu, unsigned long *capacity)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_sched_cpu_starting(void *unused, int cpu)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_sched_cpu_dying(void *unused, int cpu)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_set_task_cpu(void *unused, struct task_struct *p, unsigned int new_cpu)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_new_task_stats(void *unused, struct task_struct *p)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_account_irq(void *unused, struct task_struct *curr, int cpu, s64 delta)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_flush_task(void *unused, struct task_struct *p)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_enqueue_task(void *unused, struct rq *rq, struct task_struct *p)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_dequeue_task(void *unused, struct rq *rq, struct task_struct *p)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_update_misfit_status(void *unused, struct task_struct *p,
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_try_to_wake_up(void *unused, struct task_struct *p)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_try_to_wake_up_success(void *unused, struct task_struct *p)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_tick_entry(void *unused, struct rq *rq)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_schedule(void *unused, struct task_struct *prev,
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_resume_cpus(void *unused, struct cpumask *resuming_cpus, int *err)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_update_cpus_allowed(void *unused, struct task_struct *p,
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_sched_setaffinity(void *unused, struct task_struct *p,
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_sched_fork_init(void *unused, struct task_struct *p)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_ttwu_cond(void *unused, bool *cond)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_sched_exec(void *unused, bool *cond)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_build_perf_domains(void *unused, bool *eas_check)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_force_compatible_pre(void *unused, void *unused2)
KERNEL.PLATFORM.1.0.r1-11200-kernel.0:kernel/sched/walt/walt.c:static void android_rvh_force_compatible_post(void *unused, void *unused2)
```

Still not enough GKI is not fully enforced

- “In GKI 2.0, devices launching with kernel version 5.10 or higher must ship with the GKI kernel (beginning with Android 12).”
 - Notice that it doesn’t say “Google’s GKI kernel”
- OEMs are allowed to change the core kernel as long as it follows the GKI / KMI model

In practice

Snapdragon 8 Gen 1 (2022)

```
taro:/dev/block/bootdevice/by-name # ls
ALIGN_T0_128K_1  dtbo_b      modemst1
ALIGN_T0_128K_2  featenabler_a  modemst2
abl_a            featenabler_b  multiimgoem_a
abl_b            frp          multiimgoem_b
aop_a            fsc          multiimgqti_a
aop_b            fsg          multiimgqti_b
aop_config_a    hyp_a        mysuper_a
aop_config_b    hyp_b        persist
apdp             imagefv_a   qmcs
apdpb            imagefv_b   qupfw_a
bluetooth_a     keymaster_a  qupfw_b
bluetooth_b     keymaster_b  qweslicstore_a
boot_a           keystore     qweslicstore_b
boot_b           limits       rawdump
cdt              limits-cdsp  recovery_a
connsec          logdump     recovery_b
cpucp_a          logfs       rtice
cpucp_b          mdcompress  secdata
ddr              mdtp_a     shr_m_a
devcfg_a         mdtp_b     shr_m_b
devcfg_b         mdtpsecapp_a splash
devinfo          mdtpsecapp_b spunvm
dip              metadata    ssd
dsp_a            misc       storsec
dsp_b            modem_a   super
dtbo_a           modem_b   toolsfv
tz_a             tz_b       tzsc
tz_b             tzsc      uefi_a
uefi_b          uefi_b     uefisecapp_a
uefisecapp_b    uefisecapp_b uefivarstore
userdata         vbmeta_a   vbmeta_system_a
vbmeta_b         vbmeta_b   vbmeta_system_b
vbmota           vbmota     vendor_boot_a
vbmota           vbmota     vendor_boot_b
vm-bootsys_a    vm-bootsys_b
vm-data          xbl_a      xbl_b
xbl_a           xbl_b      xbl_config_a
xbl_b           xbl_sc_logs xbl_config_b
xbl_ramdump_a   xbl_ramdump_b xbl_sc_test_mode
xbl_ramdump_b   xbl_sc_logs xbl_sc_test_mode
```

```
taro:/ # mount | grep dm- | head -n6
/dev/block/dm-2 on / type ext4 (ro,seclabel,nodev,relatime,discard)
/dev/block/dm-3 on /system_ext type ext4 (ro,seclabel,relatime,discard)
/dev/block/dm-1 on /product type ext4 (ro,seclabel,relatime,discard)
/dev/block/dm-4 on /vendor type ext4 (ro,seclabel,relatime,discard)
/dev/block/dm-5 on /vendor_dlk type ext4 (ro,seclabel,relatime,discard)
/dev/block/dm-0 on /odm type ext4 (ro,seclabel,relatime,discard)
```

Loaded by the bootloader

boot	GKI core kernel image
vendor_boot	initramfs, first stage modules (e.g., UFS)
vendor_dlk	Rest of the kernel modules (e.g., Wi-Fi)

Loaded by the Android's init binary

In practice Snapdragon 8 Gen 1 (2022)

```
taro:/ # find /vendor_dlkm/
/vendor_dlkm/
/vendor_dlkm/etc
/vendor_dlkm/etc/build.prop
/vendor_dlkm/etc/NOTICE.xml.gz
/vendor_dlkm/etc/fs_config_dirs
/vendor_dlkm/etc/fs_config_files
/vendor_dlkm/lib
/vendor_dlkm/lib/modules
/vendor_dlkm/lib/modules/adsp_loader_dlkm.ko
...
/vendor_dlkm/lib/modules/mmrmr_test_module.ko
/vendor_dlkm/lib/modules/modules.alias
/vendor_dlkm/lib/modules/modules.blocklist
/vendor_dlkm/lib/modules/modules.dep
/vendor_dlkm/lib/modules/modules.load
/vendor_dlkm/lib/modules/modules.softdep
/vendor_dlkm/lib/modules/msm-cvp.ko
...
/vendor_dlkm/lib/modules/zsmalloc.ko
/vendor_dlkm/lost+found
```

```
# adb shell dmesg | grep 'init:'
[ 5.085137] init: init first stage started!
[ 5.090835] init: Loading module /lib/modules/qcom_hwspinlock.ko with args ''
[ 5.106962] init: Loaded kernel module /lib/modules/qcom_hwspinlock.ko
...
[ 6.914768] init: Loading module /lib/modules/msm-geni-se.ko with args ''
[ 6.929909] init: Loaded kernel module /lib/modules/msm-geni-se.ko
[ 6.936268] init: Loading module /lib/modules/msm_geni_serial.ko with args 'con_enabled=1'
[ 6.958682] init: Loaded kernel module /lib/modules/msm_geni_serial.ko
[ 6.965412] init: Loading module /lib/modules/ns.ko with args ''
[ 6.971603] init: Loaded kernel module /lib/modules/ns.ko
...
[ 7.867425] init: Loading module /lib/modules/bcl_pmic5.ko with args ''
[ 7.883855] init: Loaded kernel module /lib/modules/bcl_pmic5.ko
[ 7.890187] init: Loaded 100 kernel modules took 2800 ms
```

```
adb shell dmesg | grep 'modprobe:'
[ 13.683976] modprobe: Loading module /vendor/lib/modules/mdt_loader.ko with args ''
[ 13.702413] modprobe: Loaded kernel module /vendor/lib/modules/mdt_loader.ko
...
[ 33.188493] modprobe: Loading module /vendor/lib/modules/msm_drm.ko with args
'dsi_display0=qcom,mdss_dsi_nt36672c_fhd_plus_90hz_video:'
...
[ 41.771381] modprobe: Loaded kernel module /vendor/lib/modules/pinctrl_lpi_dlkm.ko
```

First stage (vendor_boot)

375 modules!

vendor_dlkm

In practice

Using Google-provided GKI images

The screenshot shows a web browser displaying the [Android Source documentation](https://source.android.com/docs/core/architecture/kernel/gki-android12-5_10-release-builds). The page is titled "Core Topics" and features a sidebar on the left with sections like "Getting Started", "Security", "Core Topics" (which is underlined), "Compatibility", "Android Devices", and "Reference". The main content area is titled "Android12-5.10 launch releases" and includes a bulleted list: "Parent branch: [android12-5.10](#) ([history](#))". Below this, it says "The following branches are eligible for respin, either as requested by partners or for the security patches cited in the Android Security Bulletin (ASB):". A section for "August, 2022 releases" lists a branch: "Branch: [android12-5.10-2022-08](#) ([history](#))". A table titled "Release build" provides detailed information about the release:

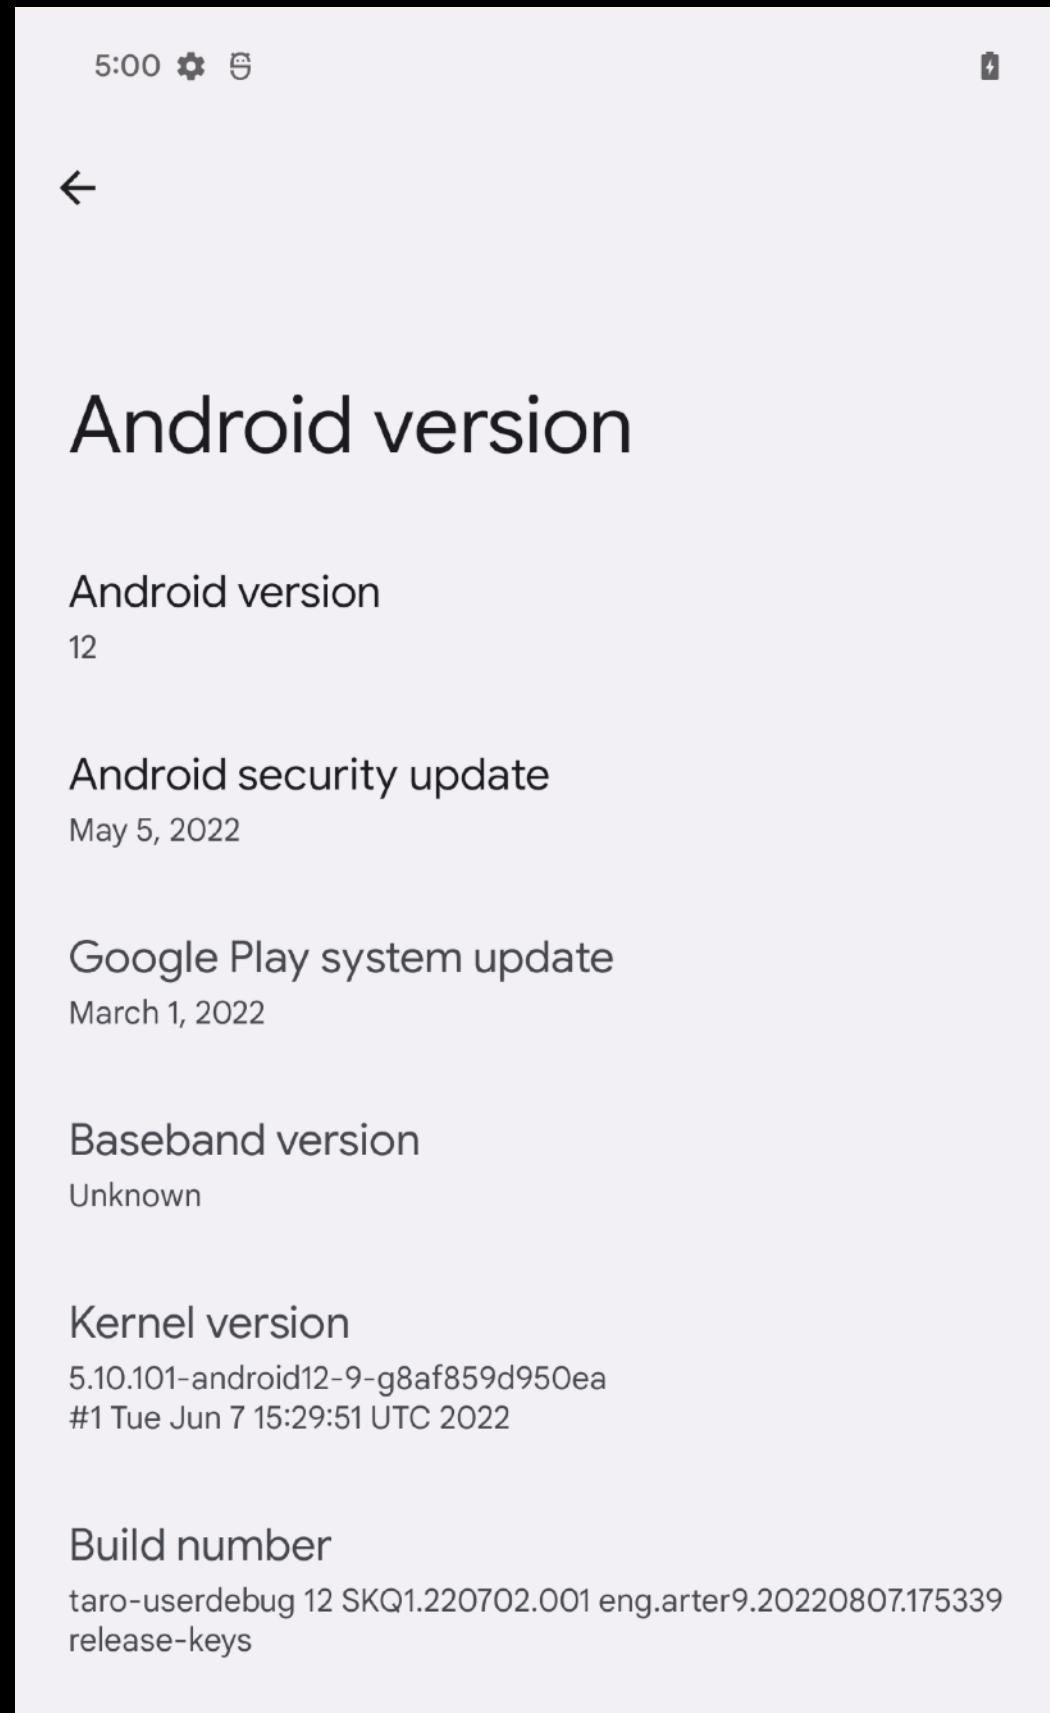
Release date	Tag	SPL	Kernel artifacts	Certified GKI
2022-08-09	android12-5.10-2022-08_r1 SHA1: 91bfc78bc009e8afc8e6	2022-09-05	kernel	boot-5.10.img boot-5.10-gz.img boot-5.10-lz4.img

The sidebar also lists "Kernel" sub-sections: Overview, Stable Releases & Updates, Android Common Kernels, The GKI project, GKI development, GKI Versioning, and "GKI Release Builds" which includes "Overview" and "Releases" for "android12-5.10" (which is highlighted).

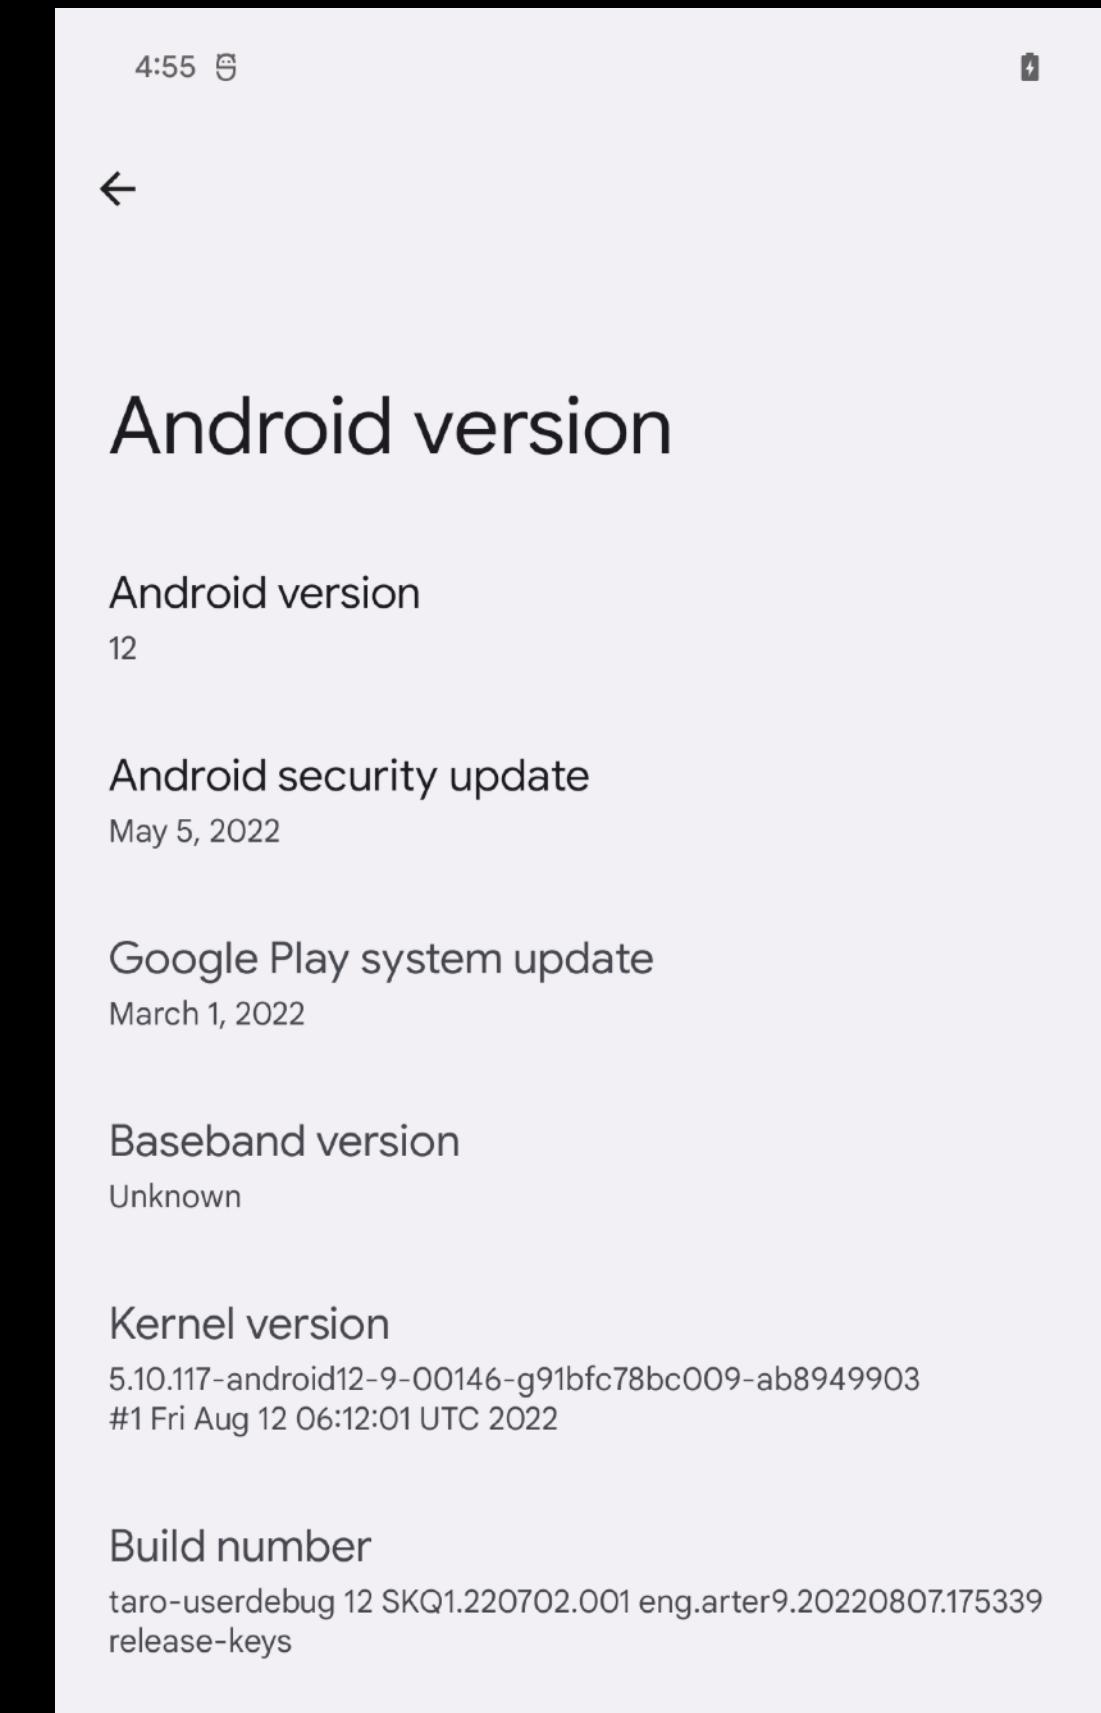
https://source.android.com/docs/core/architecture/kernel/gki-android12-5_10-release-builds

In practice

Using Google-provided GKI images



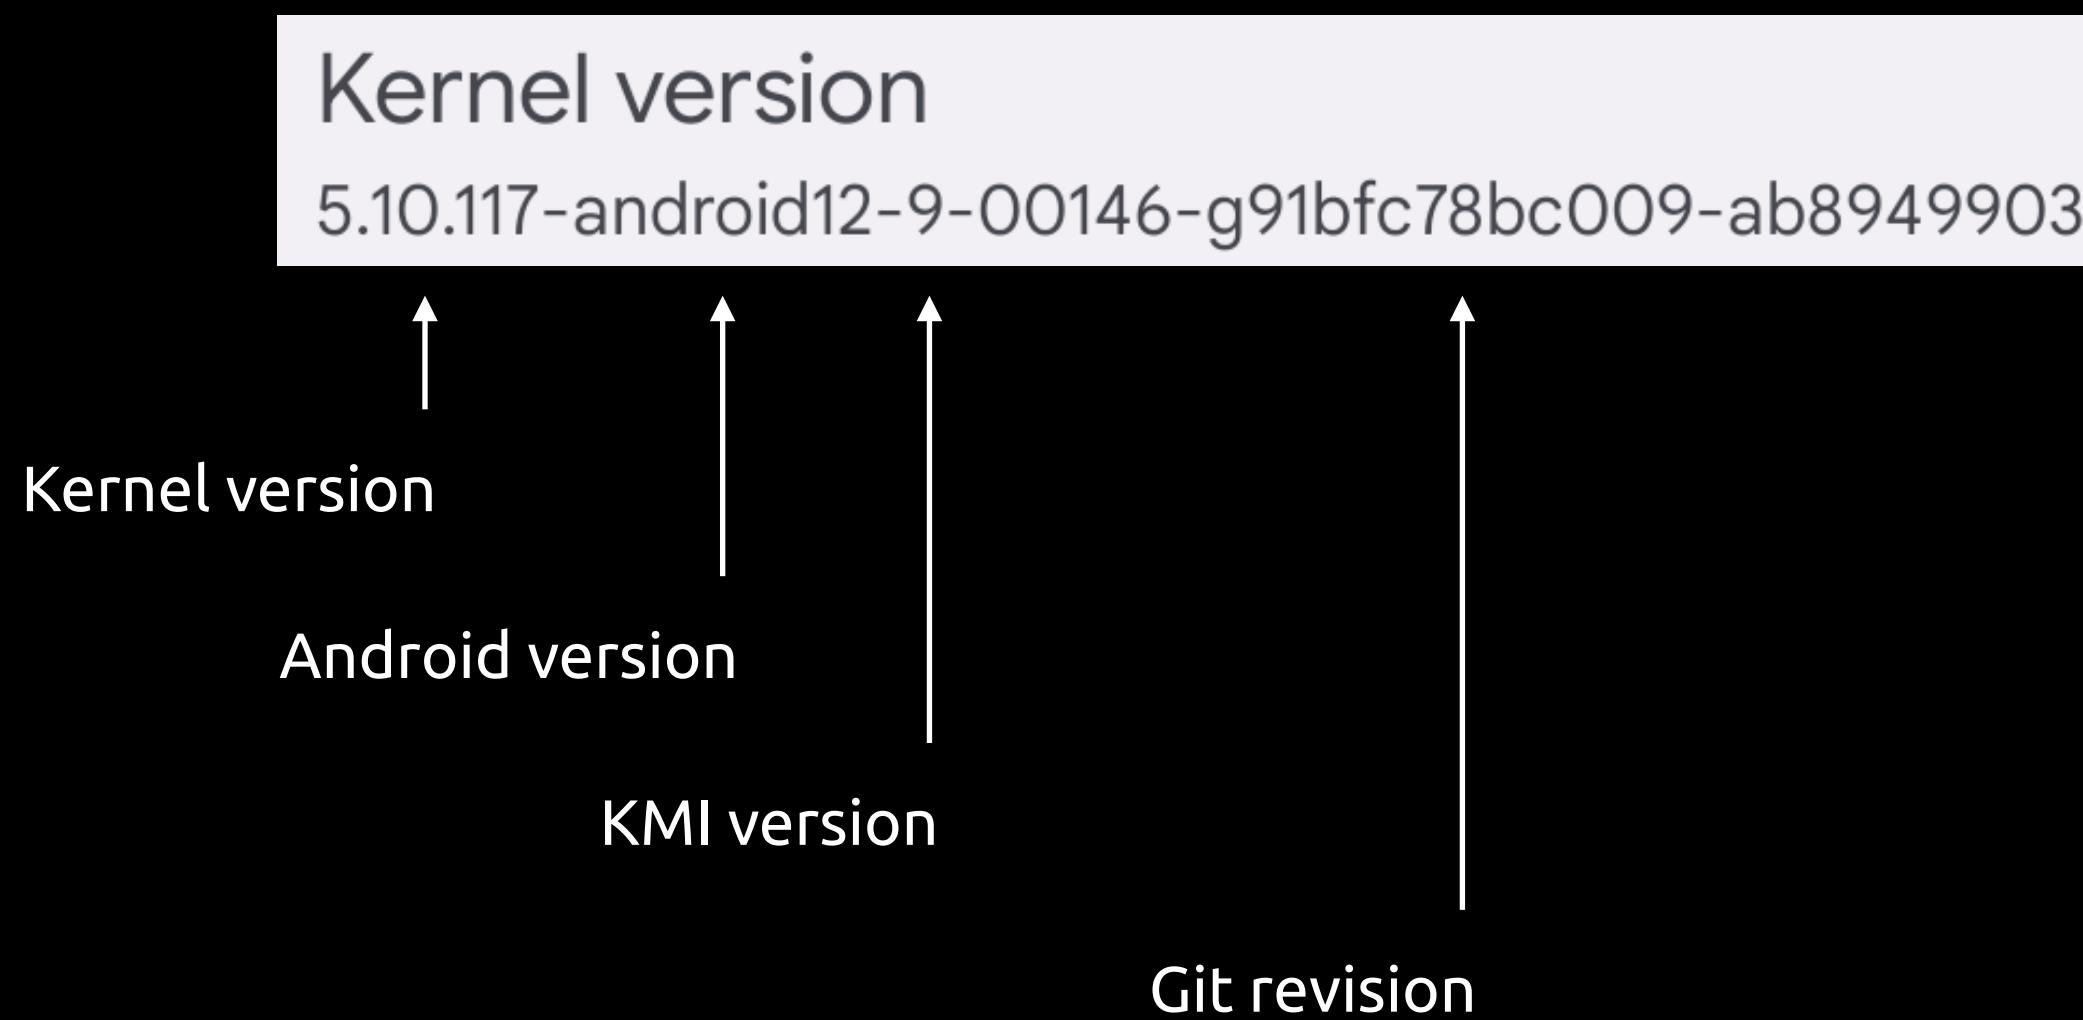
Stock kernel by Qualcomm



Google-provided GKI image

In practice

Using Google-provided GKI images



In practice

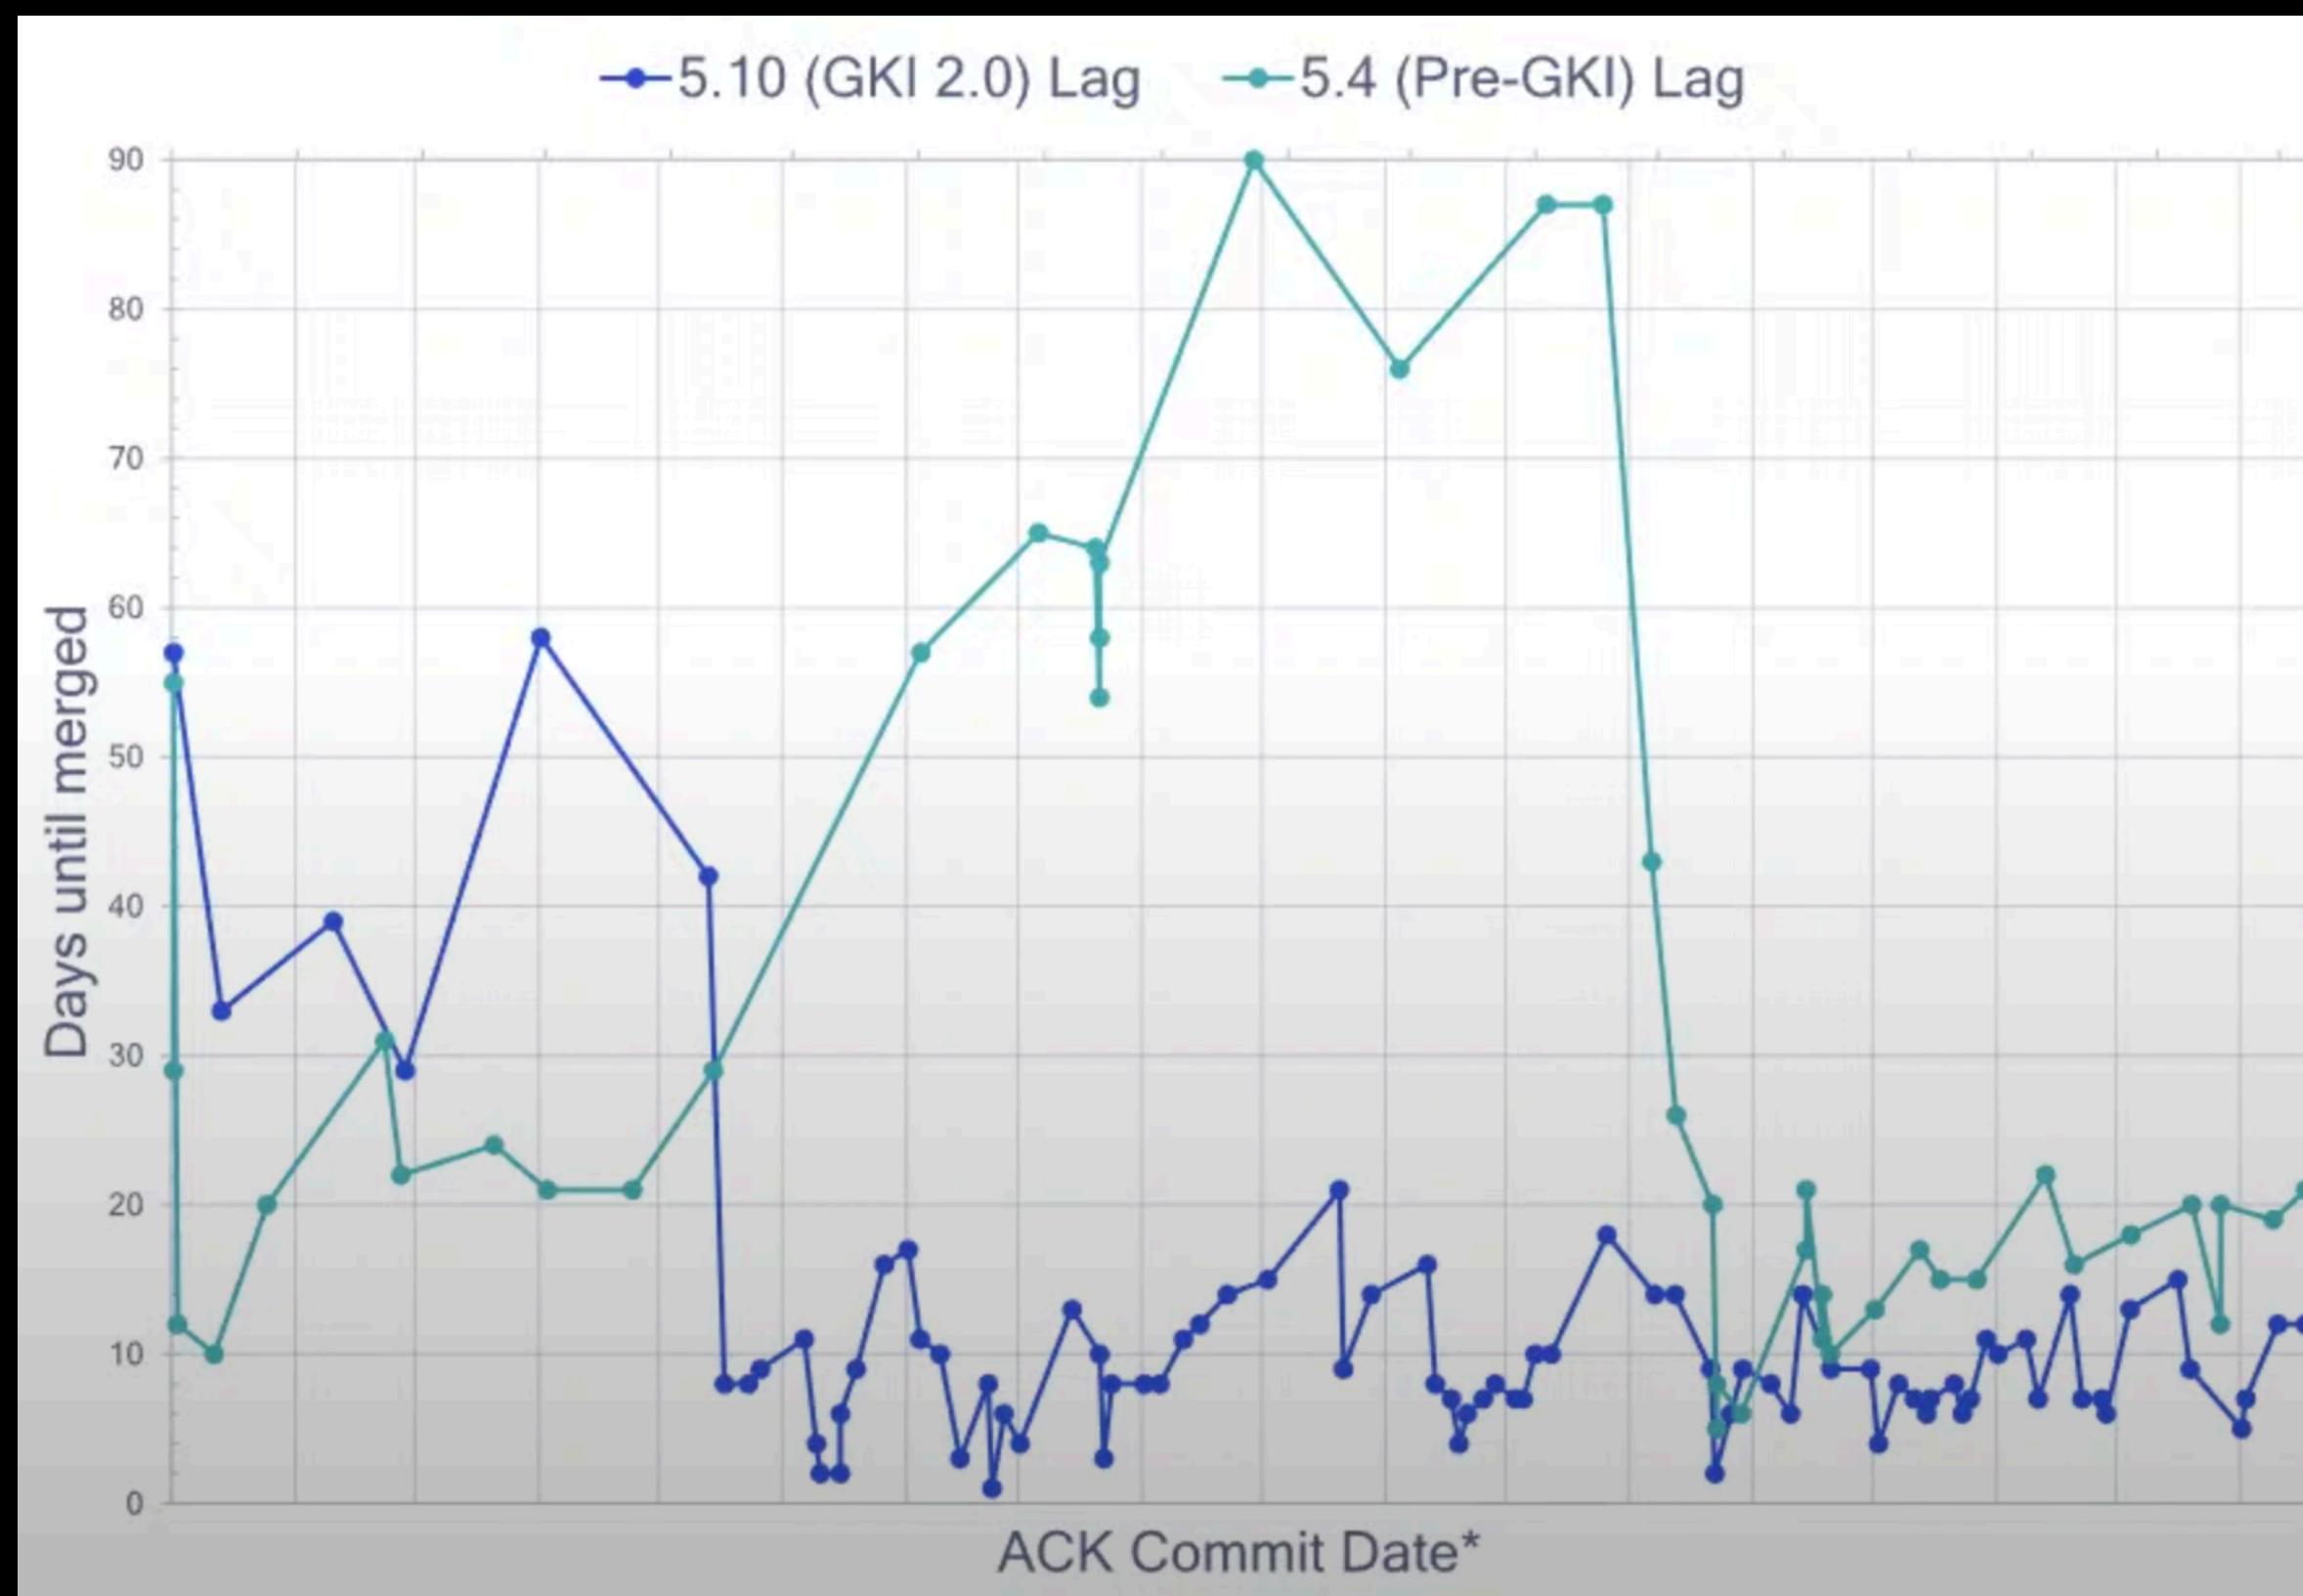
What GKI does and doesn't

- GKI allows updating core kernel and vendor modules asynchronously
- GKI guarantees module compatibilities within the same KMI version
 - 5.10.y-android12-9
 - 5.10.y-android13-5
 - 5.15.y-android13-1
- GKI **doesn't** guarantees stable ABIs across different major kernel versions
 - Upgrading a device from v5.10 to v5.15 is still a burden

In practice

How effective is it?

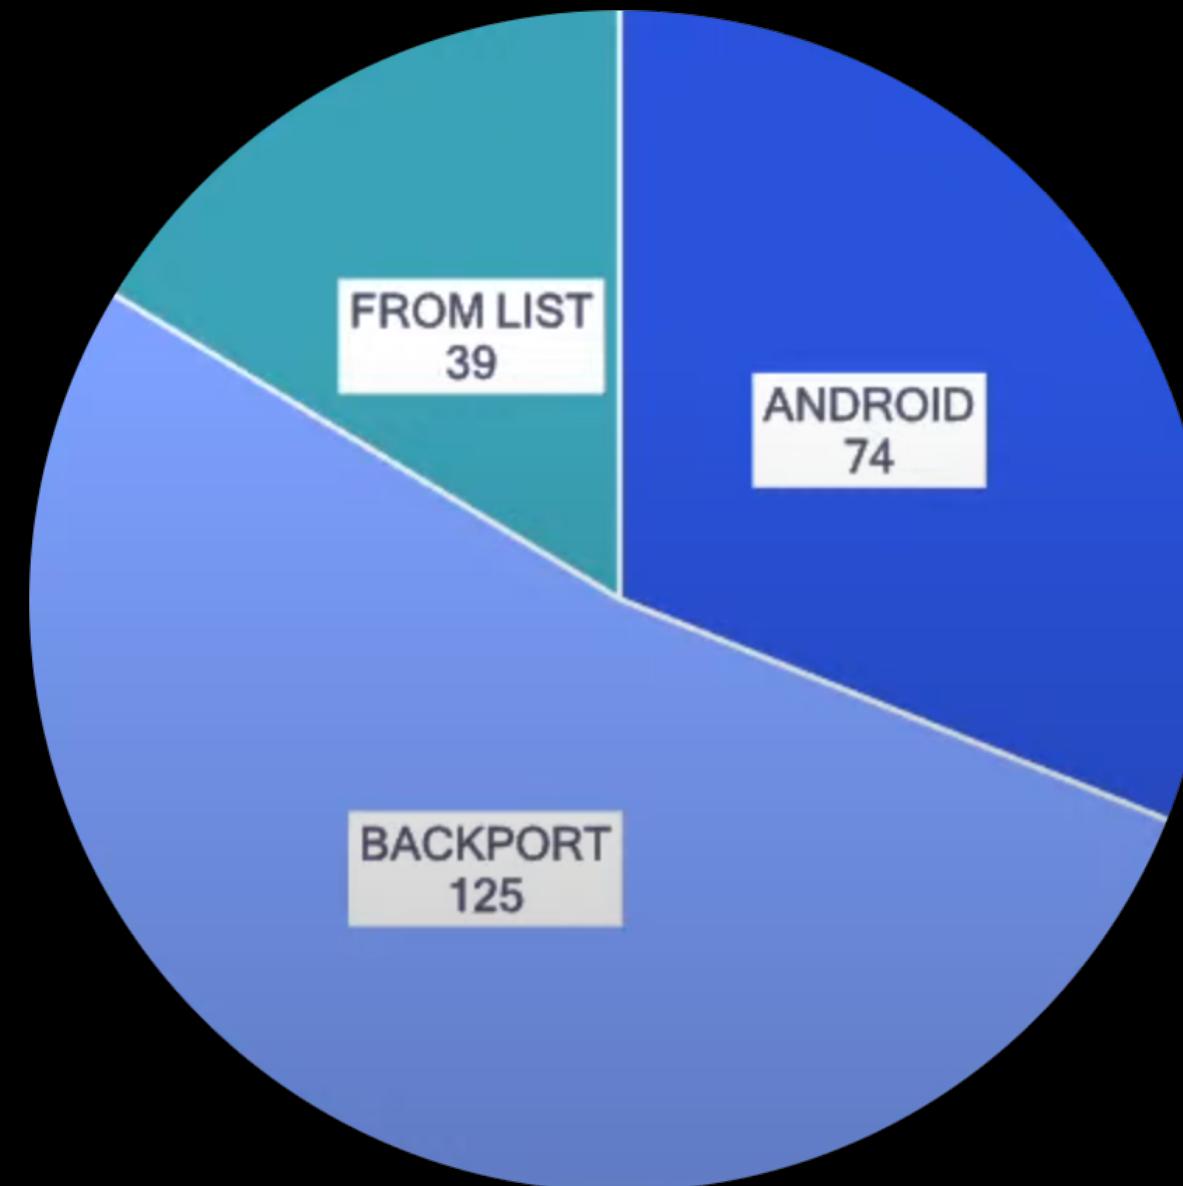
- Security patches integrated quickly without any issues (e.g., Spectre-BHB)



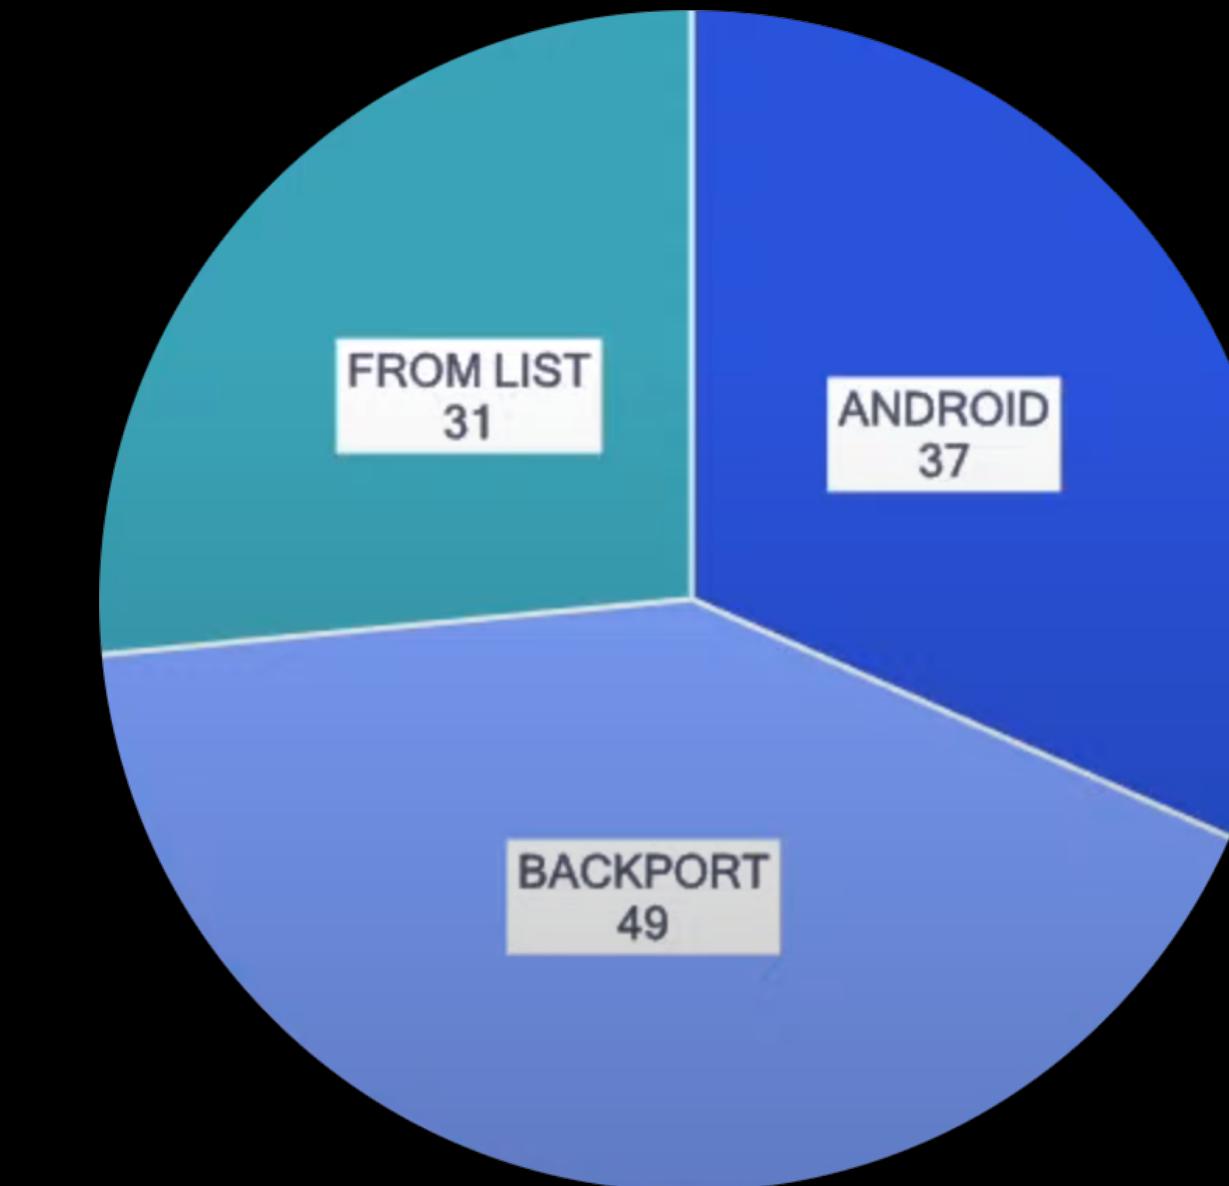
In practice

How effective is it?

android12-5.10 (238 commits)



android13-5.15 (117 commits)



GKI successfully reduces divergence from mainline

In practice

How effective is it?

- Major OEMs are not shipping Google's GKI image
 - Upstreaming every value-adds and optimizations to Google is infeasible
 - This means switching to Google's GKI later via OTA is also infeasible
 - Imagine losing features in the name of enhanced security
 - “One kernel image for all Android devices” is still left unrealized

Verdict

Current state of Android actively discourages modding

- 2011 - Close-sourced vendor modules
- 2013 - eFUSE
- 2014 - Unlocking bootloader breaks camera
- 2015 - Widevine L1
- 2017 - SafetyNet
- 2019 - R/O file-system by design, dynamic partitions
- 2021 - Re-introduction of kernel modules
- 2022 - GKI 2.0

Some things shouldn't be this hard...

Not kernel-related, but still

- Obtaining root access on unlocked bootloader
 - Magisk became a very complicated piece of software
- Replacing system files
 - Recall that system partitions are R/O by design
 - Dedup-ext4, f2fs RO, EROFS
 - overlayfs is only available on debug firmware
 - AOSP is too dependent on GMS to be usable

No one pressures OEMs for GPL

Not GKI-related, but still

- Major OEMs often take months to release the kernel's source code
 - Even never
 - Obviously without Git commit history
 - Often times broken with partial source code
 - Compartmentalized modules made things even more complicated

GKI is a nightmare to tinkerers and hackers

Build environment

- GKI now requires OEMs to build kernel and modules within the GKI build environment
 - OEMs can fork the build environment
 - Build environment itself is not covered by GPL
 - Building for Pixels is easy
 - Building for anything else is not
 - No clear instructions given alongside the source code

GKI is a nightmare to tinkerers and hackers

LTO

- Changing a single line of source code takes 20 minutes to recompile
 - On a 64-cores AMD Ryzen Threadripper 3990X
 - LTO takes too long
- Disabling LTO disables CFI
 - <https://source.android.com/docs/security/test/kcfi>
 - Disabling CFI breaks module compatibility

GKI is a nightmare to tinkerers and hackers

Partition shenanigans

- Recall that you need to change 3 partitions
 - “boot” for the core kernel image
 - “vendor_boot” for the initial kernel modules
 - “vendor_dlkm” for the rest of the kernel modules
 - “vendor_dlkm” is inside the dynamic partition
 - Your device may even lack “vendor_dlkm”, which then you need to modify the “vendor” partition directly
- You better hope you got all of it correctly done to *just boot*

GKI is a nightmare to tinkerers and hackers

Monolithic kernel

- Disabling modules will produce a monolithic kernel
 - “This will make things simpler, right?”
- OEMs no longer test initcall orders
- Various modules will simply break because of orderings
 - X module may depend on Y module to be loaded prior
 - Z module may depend on a specific userspace file to be present
 - No logs available to debug

My first experience with GKI

HDK8450

- Snapdragon 8 Gen 1 Mobile HDK
 - Linux v5.10
 - Google's GKI worked nicely
 - But I wanted full-control over the kernel
 - Thankfully, Qualcomm's environment is open-source
 - <https://git.codelinaro.org/clo/la/kernel/manifest>
 - Unfortunately, it didn't build without proprietary BSP
 - Attempts to fix it to build doesn't seem to be worth it



My first experience with GKI

HDK8450

- Monolithic kernel didn't boot
 - Expected, but still sad :(
 - Patched “vendor_boot” to enable earlycon
 - Got minicom to display kernel log!
 - A whole bunch of initcalls causing kernel panics
 - Fixing it one-by-one for all 375 modules? No thanks.



Lazy initcall

Allowing monolithic kernel

- Idea: defer initcalls until userspace asks for the corresponding module
- Mark device drivers to defer its initialization and add to the list
- Modify the `init_module(2)` system-call
 - Get the module name from the requested .ko file
 - Call the corresponding device driver's `init()` function
- <https://github.com/arter97/hdk-kernel/commit/e0085cde67a4>
- <https://github.com/arter97/hdk-kernel/commit/c425693e00d1>

Lazy initcall

Allowing monolithic kernel

```
[ 5.089408] init: init first stage started!
[ 5.095136] init: Loading module /lib/modules/qcom_hwspinlock.ko with args ''
[ 5.102852] lazy_initcall: lazy_initcalls[13]: qcom_hwspinlock's init function returned 0
[ 5.111263] init: Loaded kernel module /lib/modules/qcom_hwspinlock.ko
[ 5.117974] init: Loading module /lib/modules/smem.ko with args ''
[ 5.124460] lazy_initcall: lazy_initcalls[21]: smem's init function returned 0
[ 5.131888] init: Loaded kernel module /lib/modules/smem.ko
[ 5.137643] init: Loading module /lib/modules/minidump.ko with args ''
[ 5.146545] Minidump: Enabled with max number of regions 200
[ 5.152528] lazy_initcall: lazy_initcalls[110]: minidump's init function returned 0
[ 5.160410] init: Loaded kernel module /lib/modules/minidump.ko
```

4:55 ⚙️ 🔋



Android version

Android version

12

Android security update

May 5, 2022

Google Play system update

March 1, 2022

Baseband version

Unknown

Kernel version

5.10.110-arter97-g9a99d374727b
#1 Sun Sep 25 17:17:18 KST 2022

Build number

taro-userdebug 12 SKQ1.220702.001 eng.arter9.20220807.175339
release-keys



Easily the hardest device I ever had to work on to *just boot*

Custom build environment

- Google's GKI build environment - 20 minutes
- Full build - 50 seconds
- Full build with ccache - 17 seconds
- Incremental build - 15 seconds
- AMD Ryzen Threadripper 3990X (64-cores)

Few ideas for improvements...



- Properly mandate GPL
- Verify the released kernel source and environment properly boots and passes tests
- Verify CONFIG_MODULES=n passes tests
 - Tinkerers don't need kernel modules
 - Too many things break on monolithic kernels
- Standardize and mandate a way to retrieve kernel panic logs
 - Some are straight-up unavailable on production devices
 - pstore-blk is now available upstream

Q & A